

# EigenFace Implementation

## Basic Info

- Student ID: 3180105504
- Student Name: Xu Zhen
- Instructor Name: Song Mingli
- Course Name: Computer Vision
- Homework Name: EigenFace Implementation
- Basic Requirements:
  - Implement an EigenFace face recognition trainer/tester pair
  - Construct a face database of at least 40 persons (including the author)
  - Don't use OpenCV's API for eigenfaces
  - APIs relating to eigenvalues/eigenvectors can be used
  - No Qt for GUI, only HighGUI that comes with OpenCV can be used

I really don't know why this requirements should exist

OpenCV uses Qt as a backend to display stuff, at least it does so with modern Python interfaces

This is a screenshot when I accidentally discovered this trying to open a window with an SSH connection to my server

```
qt.qpa.xcb: could not connect to display
qt.qpa.plugin: could not load the Qt platform plugin "xcb" in "/home/xzdd/miniconda3/envs/py38/lib/python3.8/site-packages/cv2/qt/plugins" even though it was found.
This application failed to start because no Qt platform plugin could be initialized. Reinstalling the application may fix this problem.

Available platform plugins are: xcb, eglfs, minimal, minimalegl, offscreen, vnc.

[1] 153655 abort (core dumped) ./test.py image_0318.jpg model.npz.grayscale builtin.json
(py38) → homework3 git:(master) ✘ git status
```

## Experiment Principles

### Matrix & Eigenvalues & Eigenvectors

Matrix is a form of transformation (square matrices)

This is especially useful when we're trying to unveil the true nature of a lot of linear algebra concepts with the help of geometry (visualization).

Of course, one can always write abstract symbols and prove one's way throughout the linear algebra sea without having to draw even a single checkerboard, it's still gonna be extremely helpful, especially for us beginners, to understand some of the most basic concepts so deeply that, they're like imprinted into our heads.

Hereby I highly recommend the series brought to you by *3Blue1Brown*, a famous YouTuber, and a mathematician. Link below

### Essence of linear algebra

I found these masterpiece series when trying to understand eigen stuff better (eigenvalues, eigenvectors, eigenbasis and of course, eigenfaces)

The most important takeways are:

- Matrices (square) can and should be viewed as a transform where each columns represents the new basis's location

We can visualize this process by **dragging**  $[1, 0]$  to the first column of the transformation matrix, and  $[0, 1]$  to the second, assuming a 2 by 2 matrix

- Eigenvectors are vectors that doesn't get knocked off their span when performing the transformation described in the previous term

Typically, for a 2 by 2 matrix, there's two directions, which, during the **dragging** in the previous item, don't rotate around, but just scales up to a certain degree

And that **certain degree**, are the eigenvalues

OpenCV gives us some API for computing eigenvectors and eigenvalues. Unfortunately, those direct methods are a bit outdated and don't exist anymore in modern interfaces (especially so if we're using Python as the programming language for performing these CV tasks).

But as we can probably guess from the importance of eigen stuff, there's tons of other resources for us to call. And they all got optimized the heck out of them because people use them sooooo often.

## PCA: Principal Component Analysis

Intuitively, principal components are the axes, onto which if data points are projected, the projections' distances to the origin get the largest variance

Similar to eigen stuff, we find it much easier to understand PCA with a few visualization and a detailed illustration with a lot of examples.

This is the script that inspired me the most. Masterpiece indeed. As a technical report, it got over 2500 citations on Google Scholar

### A tutorial on Principal Components Analysis

The main takeways are:

- Principal Components are those axes that preserve information about the original data the most

Take a 2-d checkerboard for example, there's a bunch of points on this plane and they roughly form a straight line

For simplicity, we assume this rough line to be `y=x`

Then intuitively, if we memorize the points as the distance to the origin, we'll lose one degree of freedom

But we'll be able to roughly reconstruct the points by drawing them on the line `y=x` using their distance to the origin

And this `y=x` would be the **Principal Components** for those bunch of points

- Mathematically, one can compute the principal components by getting the eigenvectors sorted by their corresponding eigenvalues reversely of the covariance matrix of those data points

A long sentence, right?

A covariance matrix is a convenient way to write down all the covariance of the data points in question, where its index indicates the orginal data. For example, the value at `[i, j]` is the covariance of `data[i]` and `data[j]`

## Core APIs

### OpenCV API

```
1  def PCACompute(data, mean, eigenvectors= ... , maxComponents= ... ) → typing.Any:
2      """
3          PCACompute(data, mean[, eigenvectors[, maxComponents]]) → mean, eigenvectors
4          .    wrap PCA::operator()
5
6
7
8          PCACompute(data, mean, retainedVariance[, eigenvectors]) → mean, eigenvectors
9          .    wrap PCA::operator()
10         """
11         ...
```

In our implementation, we provided the user with the option of using `PCACompute` of OpenCV to compute the principal components efficiently.

Note that if we take the most basic approach, we'd have to construct a giant covariance matrix and compute the eigenvectors/eigenvalues of it.

For example, if we're to compute the covariance matrix of some 512 by 512 color image, meaning the covariance matrix would be  $3 \times 2^{18}$ : 786432 by 786432

**Consider 64-bit double for the entries, we'd be needing  $2^{36} * 9 * 2^3$  Bytes  $\rightarrow 4.5$  TiB memory to just store the matrix.**

And this would be totally unacceptable.

So we advise you to use the OpenCV's implementation for PCA, as long as you provide us with the nEigenFaces you want instead of trying to achieve some sort of **energy percentage**. This would be another issue will discuss later.

Well, but the instructions advise us to use our custom implementation of PCA, so we still implemented that version. You should use that version by tweaking the configuration file for our program.

Don't try computing covariance matrix for a 512\*512 color image set, or we'll be asking you for 4.5TiB memory.

Instead, use a smaller image for that kind of computation, which can be set through the configuration file.

Like 128\*128 gray image,  $2^{17}$  Bytes are pretty reasonable for matrix generation. And a small matrix like that CAN STILL REQUIRE TONS OF TIME TO COMPUTE EIGEN STUFF.

## Scipy API

```
1 @array_function_dispatch(_cov_dispatcher)
2 def cov(m, y=None, rowvar=True, bias=False, ddof=None, fweights=None,
3         aweights=None):
```

As mentioned above, to implement our own version of PCA computation, we'd first construct the covariance matrix for the image data using **numpy.cov**.

We're aware that there's covariance matrix API in OpenCV but the **numpy** one just comes in real handy

Then, with the covariance matrix, we're gonna have to compute its eigenvalues/eigenvectors. Here we used:

```
1 def eigsh(A, k=6, M=None, sigma=None, which='LM', v0=None,
2           ncv=None, maxiter=None, tol=0, return_eigenvectors=True,
3           Minv=None, OPinv=None, mode='normal'):
4     """
5     Find k eigenvalues and eigenvectors of the real symmetric square matrix
6     or complex hermitian matrix A.
7     """
```

to compute the eigenvalues and eigenvectors with more efficiency on a real symmetrix square matrix.

Full path of the API is

```
scipy.sparse.linalg.eigen.eigsh
```

If you're wondering why we're not using the eigenvalue/eigenvector interface of OpenCV, they don't come provided in the new interface `cv2`

## Others

If we're to retain certain amount of energy when selecting the principal components from all eigenvectors, we'll have to firstly compute all the eigenvalues and sum them up. Eventually we'll have to decide how many eigenvalues' sum (descending order) can add up to take the right proportion we want.

Yeah, we do know **the trace of a matrix** is also the sum of all eigenvalues it has. But without the actual eigenvalues, we're not able to determine **how many** eigenvalues we need (when sorted) to take up a fixed proportion of the sum.

So in the previous two sets of APIs, we gave user the privilege to compute only a fixed number of eigenvalues/eigenvectors, **which is hugely efficient compared to computing all eigenvalues for some matrix**.

```
1  def PCACompute2(data, mean, eigenvectors=..., eigenvalues=...,  
2    maxComponents=...) → typing.Any:  
3  """  
4  PCACompute2(data, mean[, eigenvectors[, eigenvalues[, maxComponents]]]) →  
5  mean, eigenvectors, eigenvalues  
6  .  wrap PCA::operator() and add eigenvalues output parameter  
7  
8  PCACompute2(data, mean, retainedVariance[, eigenvectors[, eigenvalues]]) →  
9  mean, eigenvectors, eigenvalues  
10 .  wrap PCA::operator() and add eigenvalues output parameter  
11 """  
12 ...
```

For OpenCV's implementation, we should use `PCACompute2` to compute the eigenvalues/eigenvectors efficiently. By not providing the `maxComponents` parameter, we're to get all the eigenvalues/eigenvectors.

Note that `PCACompute`, the previous interface, doesn't compute eigenvalues

## Other APIs

We used `cv2.warpAffine` to transform the input image (to match the eye with our mask)

```
1  def warpAffine(src, M, dsize=..., flags=..., borderMode=...,  
2    borderValue=...) → typing.Any:  
3  """  
4  warpAffine(src, M, dsize[, dst[, flags[, borderMode[, borderValue]]]]) → dst  
5  .  @brief Applies an affine transformation to an image.
```

```

5   .
6   .   The function warpAffine transforms the source image using the specified
7   .
8   .   \f[\texttt{dst} (x,y) = \texttt{src} (\texttt{M} _{11} x + \texttt{M} _{12} y + \texttt{M} _{13}, \texttt{M} _{21} x + \texttt{M} _{22} y + \texttt{M} _{23})\f]
9   .
10  .   when the flag #WARP_INVERSE_MAP is set. Otherwise, the transformation is
11   first inverted
12   .   with #invertAffineTransform and then put in the formula above instead of M.
13   The function cannot
14   .   operate in-place.
15   .
16   .   @param src input image.
17   .   @param dst output image that has the size dsize and the same type as src .
18   .   @param M \f$2\times 3\f$ transformation matrix.
19   .   @param dsize size of the output image.
20   .   @param flags combination of interpolation methods (see #InterpolationFlags)
21   and the optional
22   .   flag #WARP_INVERSE_MAP that means that M is the inverse transformation (
23   .   \f$\texttt{dst}\rightarrow\texttt{src}\f$ ).  

24   .   @param borderMode pixel extrapolation method (see #BorderTypes); when
25   .   borderMode=#BORDER_TRANSPARENT, it means that the pixels in the destination
26   image corresponding to
27   .   the "outliers" in the source image are not modified by the function.
28   .   @param borderValue value used in case of a constant border; by default, it
is 0.
29   .
30   .   @sa warpPerspective, resize, remap, getRectSubPix, transform
31   """
32   ...

```

And along the way we also used `cv2.getRotationMatrix2D`

```

1  def getRotationMatrix2D(center, angle, scale) → typing.Any:
2  """
3  getRotationMatrix2D(center, angle, scale) → retval
4  .   @brief Calculates an affine matrix of 2D rotation.
5  .
6  .   The function calculates the following matrix:
7  .
8  .   \f[\begin{bmatrix} \alpha & \beta & (1- \alpha ) \cdot \texttt{center.x} \\ -\beta & \alpha & \beta \cdot \texttt{center.y} \\ 0 & 0 & 1 \end{bmatrix} \f]
9  .
10 .   where
11 .
12 .   \f[\begin{array}{l} \alpha = \texttt{scale} \cdot \cos \texttt{angle} , \\ \beta = \texttt{scale} \cdot \sin \texttt{angle} \end{array}\f]
13 .
14 .

```

```

15 .   The transformation maps the rotation center to itself. If this is not the
16 target, adjust the shift.
17 .
18 .   @param center Center of the rotation in the source image.
19 .   @param angle Rotation angle in degrees. Positive values mean counter-
20 clockwise rotation (the
21 .   coordinate origin is assumed to be the top-left corner).
22 .   @param scale Isotropic scale factor.
23 .
24 """

```

We also used `cv2.equalizeHist` to equalize the histogram for grayscale image:

```

1 def equalizeHist(src, dst=... ) → typing.Any:
2 """
3     equalizeHist(src[, dst]) → dst
4     .   @brief Equalizes the histogram of a grayscale image.
5     .
6     .   The function equalizes the histogram of the input image using the following
7     algorithm:
8     .
9     .   - Calculate the histogram  $H$  for  $src$  .
10    .   - Normalize the histogram so that the sum of histogram bins is 255.
11    .   - Compute the integral of the histogram:
12    .    $H'_i = \sum_{j=0}^i H(j)$ 
13    .   - Transform the image using  $H'$  as a look-up table:  $dst(x,y) = H'(\text{src}(x,y))$ 
14    .
15    .   The algorithm normalizes the brightness and increases the contrast of the
16    image.
17    .
18    .   @param src Source 8-bit single channel image.
19    .   @param dst Destination image of the same size and type as  $src$ 
20 """

```

For color images, we shouldn't just do the histogram equalization on all RGB channels, instead we converted the image to `YCbCr` color space, then do a `histEqulization` on the Y component

**YCbCr**, **Y'CbCr**, or **Y Pb/Cb Pr/Cr**, also written as **YCBCR** or **Y'CBCR**, is a family of **color spaces** used as a part of the **color image pipeline** in **video** and **digital photography** systems. Y' is the **luma** component and CB and CR are the blue-difference and red-difference **chroma** components. Y' (with prime) is distinguished from Y, which is **luminance**, meaning that light intensity is nonlinearly encoded based on **gamma corrected RGB** primaries.

Y'CbCr color spaces are defined by a mathematical **coordinate transformation** from an associated **RGB** color space. If the underlying RGB color space is absolute, the Y'CbCr color space is an **absolute color space** as well; conversely, if the RGB space is ill-defined, so is Y'CbCr.

See [Wikipedia](#) for more.

We used `HighGUI` for image display and some other basic image IO functionality from OpenCV

We used `argparse` for constructing a user-friendly interface for our program

We used `JSON` format to store our configuration file. User can check the comment on `EigenFaceUtils.loadConfig` for a detailed illustration on the functionality of all the parameters available

We'll put it here for convenience

```
1      def loadConfig(self, filename):
2          with open(filename, "rb") as f:
3              data = json.load(f)
4          log.info(f"Loading configuration from {filename}, with content:
5              {data}")
6          self.w = data["width"] # mask width
7          self.h = data["height"] # mask height
8          self.l = np.array(data["left"]) # left eye coordinates, x, y
9          self.r = np.array(data["right"]) # right eye coordinates, x, y
10         self.isColor = data["isColor"] # whether we should treat the
11             model/input images as colored ones
12         self.nEigenFaces = data["nEigenFaces"] # target eigenfaces to construct
13
14         # setting this to null will reduce computation significantly
15         # since we'll only have to compute nEigenFaces eigenvalues/eigenvectors
16         # setting to a value would disable nEigenFaces parameter, computing
17         # from target information retain rate "targetPercentage"
18         self.targetPercentage = data["targetPercentage"]
19
20         # whether we should use cv2.PCACompute/cv2.PCACompute2 or our own PCA
21             computation
22             # note the builtin methods are much faster and memory efficient if
23             we're only requiring
24             # a small subset of all eigenvectors (don't have to allocate the
25             covariance matrix)
26             # enabling us to do large scale computation, even with colored image of
27             512 * 512
28             self.useBuiltin = data["useBuiltin"]

29             # ! deprecated
30             # Only HighGUI of OpenCV is supported.\nOther implementation removed
31             due to regulation.
32             # whether we should use matplotlib to draw results or HIGHGUI of opencv
33             # I think HIGHGUI sucks at basic figure management, I'd prefer
34             matplotlib for figure drawing
35             # but if we're looking for a more general solution for GUI, it is a
36             choice
37             self.useHighgui = data["useHighgui"]
```

And we used `numpy` and `scipy` for matrix (arrays) manipulation heavily.

Like

- `matmul`
- `norm`
- `transpose`
- `flatten`
- `expand_dims`
- `copy`
- `squeeze`
- `zeros`
- `mean`
- `min`
- `max`
- `reshape`

The model is saved using compression API for numpy

```
1  @array_function_dispatch(_savez_compressed_dispatcher)
2  def savez_compressed(file, *args, **kwds):
3      """
4          Save several arrays into a single file in compressed ``.npz`` format.
5
6          If keyword arguments are given, then filenames are taken from the keywords.
7          If arguments are passed in with no keywords, then stored filenames are
8          arr_0, arr_1, etc.
9      """
10     ...
```

We used `coloredlogs` to make our log info more intuitive for a reader and for us debugger.

## Eigenfaces

There're some basic steps to take when constructing an eigenface-based face recognizer.

We'll have to:

1. Find the eye position of a given database image

This can be done by hand (annotated by a human)

Or, for convenience we can also use a haar cascade recognizer, which just comes with OpenCV

I'm aware of the fact that you asked us not to use face recognizer to get the face of the image

But I think we're performing this experiment to **learn more about eigenfaces, not to hone our own face recognition skills by labelling all eyes by hand** when already existing eye recognizer can do things much faster. Ain't computational machines created to free us human from doing those repetitive work of computation?

We can say we're rewriting some code to **learn more**

But I don't think we can learn more **by labelling more eyes**

2. Align all the eyes to the given position of our mask.

In this process, we'll have to do some image transform:

- Translation can be determined from the difference of the mask's two eye center and the given image
- Rotation can be determined from the tangent of the vector from the left eye to the right one (or the other way around)
- Scale can be determined from length (norm) of the vector of the previous term

3. Do a **histogram equalization** on the image

For gray ones, we'd only have to perform a standard equalization, but for the colored image we'll have to firstly convert RGB to a color space including grayscale (or luminance) and perform the histogram equalization on that channel.

4. Compute the **mean face** and subtract it from all the original image data
5. Construct the **covariance matrix**. Let's flatten the aligned, equalized face to batch and compute covariance matrix of all the possible pixel indices
6. Compute the **eigenvalues/eigenvectors** of the above mentioned covariance matrix.

The eigenvectors with the largest eigenvalues are eigenfaces

7. To express the face in the form of eigenfaces, we'd take the dot product between the flattened image with all the above mentioned eigenvectors. The results that pop out are the compressed version of the image data with respect to the eigenfaces.
8. To reconstruct the face, simple to a matrix multiplication to sum up the weighted eigenvectors and add the mean face mentioned in 4. Unflatten it and we'll get the face
9. To recognize the face. Do a euclidean distance between the weights of the faces in the database and the weights computed in 8. The one with the smallest distance would be the most similar face in the database

## Implementation

---

You'll notice that our implementation is pretty long, this is because we adopted detailed logging for a better debugging experience and we supported a tons of options:

1. You can choose to treat image as grayscale or color image (reconstruction is still done to a color image)
2. You can choose to set the number of eigenfaces beforehand or just set a target information retain rate
3. You can choose the size of the mask and the eye position of the left and the right one
4. You can choose to use OpenCV's implementation of PCA or ours
5. You were also able to choose to use HighGUI or just `matplotlib`, but this part is explicitly removed due to regulation

6. Theoretically, you can use whatever dataset you want, even the eyes/face are not annotated or they're of strange size.

You don't have to stick with the dataset provided by us

Because we've done carefully alignment

7. Our implementation doesn't have to be bounded with a certain dataset

As long as it has some recognizable eye in it

(And if you provide us with the eye position, we'll polite ask you to say the as:)

```
1  """
2  We're assuming a <imageFileNameNoExt>.txt for eye position like
3  474 247 607 245
4  """
```

8. The **EigenFaceUtils** class can be configured using a configuration file

And the configuration can also be saved from a instantiated object

9. The configuration file doesn't have to have a model associated with it explicitly

The only requirements are that the mask's size are matched and **isColor** field is set correctly

So you can literally have multiple models linked to a configuration file

## Training

```
1  def train(self, path, imgext, txttext, modelName="model.npz"):
2      self.updateEyeDict(path, txttext)
3      self.updateBatchData(path, imgext)
4      if self.useBuiltIn:
5          if self.targetPercentage is not None:
6              log.info(f"Beginning builtin PCACompute2 for all
eigenvalues/eigenvectors")
7          # ! this is bad, we'll have to compute all
eigenvalues/eigenvectors to determine energy percentage
8          self.mean, self.eigenVectors, self.eigenValues =
cv2.PCACompute2(self.batch, None)
9          log.info(f"Getting eigenvalues/eigenvectors:
{self.eigenValues}, {self.eigenVectors}")
10         self.updateEigenFaces()
11         # ! dangerous, losing smaller eigenvectors (eigenvalues is
small)
12         self.eigenVectors = self.eigenVectors[0:self.nEigenFaces]
13     else:
14         log.info(f"Beginning builtin PCACompute for {self.nEigenFaces}
eigenvalues/eigenvectors")
15         self.mean, self.eigenVectors = cv2.PCACompute(self.batch, None,
maxComponents=self.nEigenFaces)
16         log.info(f"Getting mean vectorized face: {self.mean} with shape:
{self.mean.shape}")
```

```

17         log.info(f"Getting sorted eigenvectors:\n{self.eigenVectors}\nof
18             shape: {self.eigenVectors.shape}")
19     else:
20         self.updateMean()
21         self.updateCovarMatrix()
22         self.updateEigenVs()
23
24     self.updateFaceDict()
25     self.saveModel(modelName)

```

Essentially, this is the training procedure illustrated in the previous section

## Reconstruction

```

1     def reconstruct(self, img: np.ndarray) -> np.ndarray:
2         assert self.eigenVectors is not None and self.mean is not None
3
4         dst = self.unflatten(self.mean).copy() # mean face
5         flat = img.flatten().astype("float64") # loaded image with double type
6         flat = np.expand_dims(flat, 0) # viewed as 1 * (width * height *
7             color)
8         flat -= self.mean # flatten subtracted with mean face
9         flat = np.transpose(flat) # (width * height * color) * 1
10        log.info(f"Shape of eigenvectors and flat: {self.eigenVectors.shape},
11             {flat.shape}")
12
13        # nEigenFace *(width * height * color) matmul (width * height * color)
14        * 1
15        weights = np.matmul(self.eigenVectors, flat) # new data, nEigenFace *
16        1
17
18        # getting the most similar eigenface
19        eigen = self.unflatten(self.eigenVectors[np.argmax(weights)])
20
21        # getting the most similar face in the database
22        minDist = Infinity
23        minName = ""
24        for name in tqdm(self.faceDict.keys(), "Recognizing"):
25            faceWeight = self.faceDict[name]
26            dist = la.norm(weights-faceWeight)
27            # log.info(f"Getting distance: {dist}, name: {name}")
28            if dist < minDist:
29                minDist = dist
30                minName = name
31
32        log.info(f"**MOST SIMILAR FACE: {minName} WITH RESULT {minDist}**")
33        face = self.unflatten(self.mean).copy() # mean face
34        faceFlat = np.matmul(np.transpose(self.eigenVectors),
35            self.faceDict[minName]) # restored
36        faceFlat = np.transpose(faceFlat)
37        face += self.unflatten(faceFlat)

```

```

33
34         # luckily, transpose of eigenvector is its inversion
35         # Eigenvectors of real symmetric matrices are orthogonal
36         # ! the magic happens here
37         # data has been lost because nEigenFaces is much smaller than the image
38             dimension span
39             # which is width * height * color
40             # but because we're using PCA (principal components), most of the
41             information will still be retained
42             flat = np.matmul(np.transpose(self.eigenVectors), weights) # restored
43             log.info(f"Shape of flat: {flat.shape}")
44             flat = np.transpose(flat)
45             dst += self.unflatten(flat)
46             if self.isColor:
47                 ori = np.zeros((self.h, self.w, 3))
48             else:
49                 ori = np.zeros((self.h, self.w))
50             try:
51                 ori = self.getImage(minName)
52                 log.info(f"Successfully loaded the original image: {minName}")
53             except FileNotFoundError as e:
54                 log.error(e)
55             dst = self.normalizeImg(dst)
56             eigen = self.normalizeImg(eigen)
57             face = self.normalizeImg(face)
58             return dst, eigen, face, ori, minName

```

The reconstruction procedure are also already illustrated.

## Modules

### Face Alignment

Align all the eyes to the given position of our mask.

In this process, we'll have to do some image transform:

- Translation can be determined from the difference of the mask's two eye center and the given image
- Rotation can be determined from the tangent of the vector from the left eye to the right one (or the other way around)
- Scale can be determined from length (norm) of the vector of the previous term

```

1     def alignFace2Mask(self, face: np.ndarray, left: np.ndarray, right:
2         np.ndarray) → np.ndarray:
3         faceVect = left - right
4         maskVect = self.l - self.r
5         log.info(f"Getting faceVect: {faceVect} and maskVect: {maskVect}")
6         faceNorm = np.linalg.norm(faceVect)
7         maskNorm = np.linalg.norm(maskVect)
8         log.info(f"Getting faceNorm: {faceNorm} and maskNorm: {maskNorm}")
9         scale = maskNorm / faceNorm

```

```

9         log.info(f"Should scale the image to: {scale}")
10        faceAngle = np.degrees(np.arctan2(*faceVect))
11        maskAngle = np.degrees(np.arctan2(*maskVect))
12        angle = maskAngle - faceAngle
13        log.info(f"Should rotate the image: {maskAngle} - {faceAngle} = {angle}
degrees")
14        faceCenter = (left+right)/2
15        maskCenter = (self.l+self.r) / 2
16        log.info(f"Getting faceCenter: {faceCenter} and maskCenter:
{maskCenter}")
17        translation = maskCenter - faceCenter
18        log.info(f"Should translate the image using: {translation}")
19
20        if scale > 1:
21            # if we're scaling up, we should first translate then do the
scaling
22            # else the image will get cropped
23            # and we'd all want to use the larger destination width*height
24            M = np.array([[1, 0, translation[0]],
25                          [0, 1, translation[1]]])
26            face = cv2.warpAffine(face, M, (self.w, self.h))
27            M = cv2.getRotationMatrix2D(tuple(maskCenter), angle, scale)
28            face = cv2.warpAffine(face, M, (self.w, self.h))
29        else:
30            # if we're scaling down, we should first rotate and scale then
translate
31            # else the image will get cropped
32            # and we'd all want to use the larger destination width*height
33            M = cv2.getRotationMatrix2D(tuple(faceCenter), angle, scale)
34            face = cv2.warpAffine(face, M, (face.shape[1], face.shape[0]))
35            M = np.array([[1, 0, translation[0]],
36                          [0, 1, translation[1]]])
37            face = cv2.warpAffine(face, M, (self.w, self.h))
38        return face

```

## Histogram Equalization

Do a **histogram equalization** on the image

For gray ones, we'd only have to perform a standard equalization, but for the colored image we'll have to firstly convert RGB to a color space including grayscale (or luminance) and perform the histogram equalization on that channel.

```

1     @staticmethod
2     def equalizeHistColor(img):
3         ycrcb = cv2.cvtColor(img, cv2.COLOR_BGR2YCR_CB)
4         channels = cv2.split(ycrcb)
5         log.info(f"Getting # of channels: {len(channels)}")
6         cv2.equalizeHist(channels[0], channels[0])
7         cv2.merge(channels, ycrcb)
8         cv2.cvtColor(ycrcb, cv2.COLOR_YCR_CB2BGR, img)
9         return img

```

## Automate the above procedure

```
1      def getImage(self, name, manual_check=False) → np.ndarray:
2          # the load the image accordingly
3          if self.isColor:
4              img = cv2.imread(name, cv2.IMREAD_COLOR)
5          else:
6              img = cv2.imread(name, cv2.IMREAD_GRAYSCALE)
7
8          # try getting eye position
9          eyes = self.getEyes(name, img)
10         log.info(f"Getting eyes: {eyes}")
11         if not len(eyes) == 2:
12             log.warning(f"Cannot get two eyes from this image: {name},
13 {len(eyes)} eyes")
14             raise EigenFaceException("Bad image")
15
16         # align according to eye position
17         dst = self.alignFace2Mask(img, eyes[0], eyes[1])
18
19         # hist equalization
20         if self.isColor:
21             dst = self.equalizeHistColor(dst)
22         else:
23             dst = cv2.equalizeHist(dst)
24
25         # should we check every image before/after loading?
26         if manual_check:
27             cv2.imshow(name, dst)
28             cv2.waitKey()
29             cv2.destroyWindow(name)
30         return dst
31
32     def getImageFull(self, imgname) → np.ndarray:
33         # Load the image specified and check for corresponding txt file to get
34         # the eye position from the file
35         txtname = f"{os.path.splitext(imgname)[0]}.txt"
36         if os.path.isfile(txtname):
37             self.updateEyeDictEntry(txtname)
38
39         log.info(f"Loading image: {imgname}")
40         return self.getImage(imgname)
41
41     def updateBatchData(self, path=("./", ext=".jpg", manual_check=False,
42     append=False) → np.ndarray:
43         # get all image from a path with a specific extension
44         # align them, update histogram and add the to self.batch
45         # adjust logging level to be quite or not
46         prevLevel = coloredlogs.get_level()
46         if not manual_check:
```

```

47             coloredlogs.set_level("WARNING")
48
49         self.pathList = os.listdir(path)
50         self.pathList = [os.path.join(path, name) for name in self.pathList if
51 name.endswith(ext)]
52         names = self.pathList
53         if not append:
54             if self.isColor:
55                 self.batch = np.ndarray((0, self.colorLen)) # assuming color
56             else:
57                 self.batch = np.ndarray((0, self.grayLen))
58         bads = []
59         for index, name in tqdm(enumerate(names), desc="Processing batch"):
60             try:
61                 dst = self.getImage(name, manual_check)
62                 flat = dst.flatten()
63                 flat = np.reshape(flat, (1, len(flat)))
64                 self.batch = np.concatenate([self.batch, flat])
65             except EigenFaceException as e:
66                 log.warning(e)
67                 bads.append(index)
68         for bad in bads[::-1]:
69             del names[bad]
70
71         coloredlogs.set_level(prevLevel)
72         log.info(f"Getting {len(names)} names and {self.batch.shape[0]} "
73                  "batch")
74         return self.batch
75
76     def updateEyeDict(self, path=".\"", ext=".eye", manual_check=False) →
77         dict:
78             # get all possible eyes position from the files with a specific
79             # extension
80             # and add the to self.eyeDict
81             prevLevel = coloredlogs.get_level()
82             if not manual_check:
83                 coloredlogs.set_level("WARNING")
84
85             names = os.listdir(path)
86             names = [os.path.join(path, name) for name in names if
87 name.endswith(ext)]
88             log.info(f"Good names: {names}")
89             for name in names:
90                 # iterate through all txt files
91                 self.updateEyeDictEntry(name)
92
93             # restore the logging level
94             coloredlogs.set_level(prevLevel)
95             return self.eyeDict
96
97     def updateEyeDictEntry(self, name):
98         # update the dictionary but only one entry

```

```

94         with open(name, "r") as f:
95             lines = f.readlines()
96             log.info(f"Processing: {name}")
97             for line in lines: # actually there should only be one line
98                 line = line.strip() # get rid of starting/ending space \n
99                 # assuming # starting line to be comment
100                if line.startswith("#"): # get rid of comment file
101                    log.info(f"Getting comment line: {line}")
102                    continue
103                coords = line.split()
104                name = os.path.basename(name) # get file name
105                name = os.path.splitext(name)[0] # without ext
106                if len(coords) == 4:
107                    self.eyeDict[name] =
108                        np.reshape(np.array(coords).astype(int), [2, 2])
109                        order = np.argsort(self.eyeDict[name][:, 0]) # sort by
110                        first column, which is x
111                        self.eyeDict[name] = self.eyeDict[name][order]
112                    else:
113                        log.error(f"Wrong format for file: {name}, at line:
114                            {line}")
115
116            return self.eyeDict[name]

```

## Get Mean Face

Compute the **mean face** and subtract it from all the original image data

```

1     def updateMean(self):
2         assert self.batch is not None
3         # get the mean values of all the vectorized faces
4         self.mean = np.reshape(np.mean(self.batch, 0), (1, -1))
5         log.info(f"Getting mean vectorized face: {self.mean} with shape:
6             {self.mean.shape}")
7
8         return self.mean

```

## Get Covariance Matrix (if needed)

Construct the **covariance matrix**. Let's flatten the aligned, equalized face to batch and compute covariance matrix of all the possible pixel indices

This procedure would not be necessary if we want to use the OpenCV's implementation of **PCACompute**

```

1     def updateCovarMatrix(self) → np.ndarray:
2         assert self.batch is not None and self.mean is not None
3         log.info(f"Trying to compute the covariance matrix")
4         # covariance matrix of all the pixel location: width * height * color
5         self.covar = np.cov(np.transpose(self.batch-self.mean)) # subtract mean
6         log.info(f"Getting covar of shape: {self.covar.shape}")
7         log.info(f"Getting covariance matrix:\n{self.covar}")
8
9         return self.covar

```

A custom slow implementation:

```
1     def updateCovarMatrixSlow(self) → np.ndarray:
2         assert self.batch is not None, "Should get sample batch before
3             computing covariance matrix"
4         nSamples = self.batch.shape[0]
5         self.covar = np.zeros((nSamples, nSamples))
6         for k in tqdm(range(nSamples**2), "Getting covariance matrix"):
7             i = k // nSamples
8             j = k % nSamples
9             linei = self.batch[i]
10            linej = self.batch[j]
11            # naive!!!
12            if self.covar[j][i] != 0:
13                self.covar[i][j] = self.covar[j][i]
14            else:
15                self.covar[i][j] = self.getCovar(linei, linej)
16
17    @staticmethod
18    def getCovar(linei, linej) → np.ndarray:
19        # naive
20        meani = np.mean(linei)
21        meanj = np.mean(linej)
22        unbiasedi = linei - meani
23        unbiasedj = linej - meanj
24        multi = np.dot(unbiasedi, unbiasedj)
25        multi /= len(linei) - 1
26        return multi
```

## Compute Eigenvalues/Eigenvectors

Compute the **eigenvalues/eigenvectors** of the above mentioned covariance matrix.

The eigenvectors with the largest eigenvalues are eigenfaces

```
1     def updateEigenVs(self) → np.ndarray:
2         assert self.covar is not None
3
4         if self.targetPercentage is not None:
5             log.info(f"Begin computing all eigenvalues")
6             self.eigenValues = la.eigvalsh(self.covar)
7             self.eigenValues = np.sort(self.eigenValues)[::-1] # this should
8             be sorted
9             log.info(f"Getting all eigenvalues:\n{self.eigenValues}\nof shape:
10 {self.eigenValues.shape}")
11             self.updateEigenFaces()
12
13         log.info(f"Begin computing {self.nEigenFaces}
14 eigenvalues/eigenvectors")
```

```

12         self.eigenValues, self.eigenVectors = sla.eigen.eigsh(self.covar,
13 k=self.nEigenFaces)
14         log.info(f"Getting {self.nEigenFaces} eigenvalues and eigenvectors with
15 shape {self.eigenVectors.shape}")
16         # always needed right?
17         self.eigenVectors = np.transpose(self.eigenVectors.astype("float64"))
18
19         # ? probably not necessary?
20         # might already be sorted according to la.eigen.eigs' algorithm
21         order = np.argsort(self.eigenValues)[::-1]
22         self.eigenValues = self.eigenValues[order]
23         self.eigenVectors = self.eigenVectors[order]
24
25         log.info(f"Getting sorted eigenvalues:\n{self.eigenValues}\nof shape:
26 {self.eigenValues.shape}")
27         log.info(f"Getting sorted eigenvectors:\n{self.eigenVectors}\nof shape:
28 {self.eigenVectors.shape}")

```

## Face Dict

And of course we'd like to express the database with our new eigenface basis

```

1     def updateFaceDict(self) → dict:
2         # compute the face dictionary
3         assert self.pathList is not None and self.batch is not None and
4             self.eigenVectors is not None
5         # note that names and vectors in self.batch are linked through index
6         assert len(self.pathList) == self.batch.shape[0], f"
7             {len(self.pathList)} ≠ {self.batch.shape[0]}"
8         for index in tqdm(range(len(self.pathList)), "FaceDict"):
9             name = self.pathList[index]
10            flat = self.batch[index]
11            flat = np.expand_dims(flat, 0) # viewed as 1 * (width * height *
12            color)
13            flat = np.transpose(flat) # (width * height * color) * 1
14            # log.info(f"Shape of eigenvectors and flat:
15            # {self.eigenVectors.shape}, {flat.shape}")
16
17            # nEigenFace *(width * height * color) matmul (width * height *
18            color) * 1
19            weights = np.matmul(self.eigenVectors, flat) # new data,
20            nEigenFace * 1
21            self.faceDict[name] = weights
22
23            log.info(f"Got face dict of length {len(self.faceDict)}")
24
25            return self.faceDict

```

## Utilities

### Load/Save Configuration File

```
1      @staticmethod
2      def randcolor():
3          ...
4          Generate a random color, as list
5          ...
6          return [random.randint(0, 256) for _ in range(3)]
7
8      def loadConfig(self, filename):
9          with open(filename, "rb") as f:
10              data = json.load(f)
11              log.info(f"Loading configuration from {filename}, with content:
12 {data}")
13              self.w = data["width"] # mask width
14              self.h = data["height"] # mask height
15              self.l = np.array(data["left"]) # left eye coordinates, x, y
16              self.r = np.array(data["right"]) # right eye coordinates, x, y
17              self.isColor = data["isColor"] # whether we should treat the
model/input images as colored ones
18              self.nEigenFaces = data["nEigenFaces"] # target eigenfaces to construct
19
20              # setting this to null will reduce computation significantly
21              # since we'll only have to compute nEigenFaces eigenvalues/eigenvectors
22              # setting to a value would disable nEigenFaces parameter, computing
from target information retain rate "targetPercentage"
23              self.targetPercentage = data["targetPercentage"]
24
25              # whether we should use cv2.PCACompute/cv2.PCACompute2 or our own PCA
computation
26              # note the builtin methods are much faster and memory efficient if
we're only requiring
27              # a small subset of all eigenvectors (don't have to allocate the
covariance matrix)
28              # enabling us to do large scale computation, even with colored image of
512 * 512
29              self.useBuiltin = data["useBuiltin"]
30
31              # ! depricated
32              # Only HighGUI of OpenCV is supported.\nOther implementation removed
due to regulation.
33              # whether we should use matplotlib to draw results or HIGHGUI of opencv
34              # I think HIGHGUI sucks at basic figure management, I'd prefer
matplotlib for figure drawing
35              # but if we're looking for a more general solution for GUI, it is a
choice
36              self.useHighgui = data["useHighgui"]
37      def saveConfig(self, filename):
```

```

38         data = {}
39         data["width"] = self.w
40         data["height"] = self.h
41         data["left"] = self.l.tolist()
42         data["right"] = self.r.tolist()
43         data["isColor"] = self.isColor
44         data["nEigenFaces"] = self.nEigenFaces
45         data["targetPercentage"] = self.targetPercentage
46         data["useBuiltin"] = self.useBuiltin
47         data["useHighgui"] = self.useHighgui
48
49         log.info(f"Dumping configuration to {filename}, with content: {data}")
50         with open(filename, "wb") as f:
51             json.dump(data, f)

```

## Load/Save Eigenmodel

```

1     def loadModel(self, modelName):
2         # load previous eigenvectors/mean value
3         log.info(f"Loading model: {modelName}")
4         data = np.load(modelName, allow_pickle=True)
5         try:
6             self.eigenVectors = data["arr_0"]
7         except KeyError as e:
8             log.error(f"Cannot load eigenvectors, {e}")
9         try:
10            self.mean = data["arr_1"]
11        except KeyError as e:
12            log.error(f"Cannot load mean face data, {e}")
13        try:
14            self.faceDict = data["arr_2"].item()
15        except KeyError as e:
16            log.error(f"Cannot load face dict, {e}")
17        log.info(f"Getting mean vectorized face: {self.mean} with shape:
{self.mean.shape}")
18        log.info(f"Getting sorted eigenvectors:\n{self.eigenVectors}\nof shape:
{self.eigenVectors.shape}")
19        log.info(f"Getting face dict of length: {len(self.faceDict)}")
20
21    def saveModel(self, modelName):
22        log.info(f"Saving model: {modelName}")
23        np.savez_compressed(modelName, self.eigenVectors, self.mean,
24        self.faceDict)
24        log.info(f"Model: {modelName} saved")

```

## Others

And, there also some other utility functions

```
1      def unflatten(self, flat: np.ndarray) → np.ndarray:
2          # robust method for reverting a flat matrix
3          if len(flat.shape) == 2:
4              length = flat.shape[1]
5          else:
6              length = flat.shape[0]
7          if length == self.grayLen:
8              if self.isColor:
9                  log.warning("You're reshaping a grayscale image when color is
10 wanted")
11             return np.reshape(flat, (self.h, self.w))
12         elif length == self.colorLen:
13             if not self.isColor:
14                 log.warning("You're reshaping a color image when grayscale is
15 wanted")
16             return np.reshape(flat, (self.h, self.w, 3))
17         else:
18             raise EigenFaceException(f"Unsupported flat array of length:
{length}, should provide {self.grayLen} or {self.colorLen}")
19
20     def uint8unflatten(self, flat):
21         # for displaying
22         img = self.unflatten(flat)
23         return img.astype("uint8")
24
25     def updatenEigenFaces(self) → int:
26         assert self.eigenValues is not None
27         # get energy
28         self.nEigenFaces = len(self.eigenValues)
29         targetValue = np.sum(self.eigenValues) * self.targetPercentage
30         accumulation = 0
31         for index, value in enumerate(self.eigenValues):
32             accumulation += value
33             if accumulation > targetValue:
34                 self.nEigenFaces = index + 1
35                 log.info(f"For a energy percentage of {self.targetPercentage},
we need {self.nEigenFaces} vectors from {len(self.eigenValues)}")
36                 break # current index should be nEigenFaces
37
38
39     # ! unused
40     def updateEigenFaces(self) → np.ndarray:
41         assert self.eigenVectors is not None
42         log.info(f"Computing eigenfaces")
```

```

43         self.eigenFaces = np.array([self.unflatten(vector) for vector in
44             self.eigenVectors])
45         log.info(f"Getting eigenfaces of shape {self.eigenFaces.shape}")
46         return self.eigenFaces
47
48     @staticmethod
49     def normalizeImg(mean: np.ndarray) → np.ndarray:
50         return ((mean-np.min(mean)) * 255 / (np.max(mean)-
51             np.min(mean))).astype("uint8")
52
53     def drawEigenFaces(self, rows=None, cols=None) → np.ndarray:
54         assert self.eigenFaces is not None
55         # get a canvas for previewing the eigenfaces
56         faces = self.eigenFaces
57         if rows is None or cols is None:
58             faceCount = faces.shape[0]
59             rows = int(math.sqrt(faceCount)) # truncating
60             cols = faceCount // rows
61             while rows*cols < faceCount:
62                 cols += 1
63                 # has to be enough
64             else:
65                 faceCount = rows * cols
66             log.info(f"Getting canvas for rows: {rows}, cols: {cols}, faceCount:
67             {faceCount}")
68
69             if self.isColor:
70                 canvas = np.zeros((rows * faces.shape[1], cols * faces.shape[2],
71                     3), dtype="uint8")
72             else:
73                 canvas = np.zeros((rows * faces.shape[1], cols * faces.shape[2]),
74                     dtype="uint8")
75
76             for index in range(faceCount):
77                 i = index // cols
78                 j = index % cols
79                 canvas[i * faces.shape[1]:(i+1)*faces.shape[1], j * faces.shape[2]:(j+1)*faces.shape[2]] = self.normalizeImg(faces[index])
80                 log.info(f"Filling EigenFace of {index} at {i}, {j}")
81
82             return canvas
83
84     def getMeanFace(self) → np.ndarray:
85         assert self.eigenFaces is not None
86         faces = self.eigenFaces
87         mean = np.mean(faces, 0)
88         mean = np.squeeze(mean)
89         mean = self.normalizeImg(mean)
90         log.info(f"Getting mean eigenface\n{mean}\nof shape: {mean.shape}")
91         return mean

```

# Usage and Experiments

---

Check out the `train.py` and `test.py` file for more detailed implementation

And there's another list of options:

You'll notice that our implementation is pretty long, this is because we adopted detailed logging for a better debugging experience and we supported a tons of options:

1. You can choose to treat image as grayscale or color image (reconstruction is still done to a color image)
2. You can choose to set the number of eigenfaces beforehand or just set a target information retain rate
3. You can choose the size of the mask and the eye position of the left and the right one
4. You can choose to use OpenCV's implementation of PCA or ours
5. You were also able to choose to use HighGUI or just `matplotlib`, but this part is explicitly removed due to regulation
6. Theoretically, you can use whatever dataset you want, even the eyes/face are not annotated or they're of strange size.

`You don't have to stick with the dataset provided by us`

Because we've done carefully alignment

7. Our implementation doesn't have to be bounded with a certain dataset

As long as it has some recognizable eye in it

(And if you provide us with the eye position, we'll polite ask you to say the as:)

```
1  """
2  We're assuming a <imageFileNameNoExt>.txt for eye position like
3  474 247 607 245
4  """
```

8. The `EigenFaceUtils` class can be configured using a configuration file

And the configuration can also be saved from a instantiated object

9. The configuration file doesn't have to have a model associated with it explicitly

The only requirements are that the mask's size are matched and `isColor` field is set correctly

So you can literally have multiple models linked to a configuration file

## Basic Usage

Basically, you set your configuration file like this:

```
1  {
2      "width": 512,
3      "height": 512,
```

```
4         "left": [
5             188,
6             188
7         ],
8         "right": [
9             324,
10            188
11        ],
12         "isColor": true,
13         "nEigenFaces": 1000,
14         "targetPercentage": null,
15         "useBuiltIn": true,
16         "useHighgui": true
17     }
```

Eye positions would be pretty straight forward.

**isColor** defines how our program will treat the input images and generate output image

**nEigenFaces** defines the number of eigenfaces you want to generate, this option can be overwritten by **targetPercentage**, which is the target information retain rate you want to achieve.

If `targetPercentage` is set to `null` here (`None` in python implementation), we'll use `nEigenFaces`. Else we'll try computing `nEigenFaces` manually.

`useBuiltIn` means you want to use the `PCACompute` implementation of OpenCV, which is way more faster and memory efficient than our own version.

**useHighgui** is deprecated. Leave it to **true**

Then you can run the `train.py` or `test.py` with the correct arguments provided:

You can print the help message using `python train.py -h` or `python test.py -h`

```

16          The text file extension we want to read eye positions
17          off
18          from, usually .txt. But others will work too
19      -c CONFIG, --config CONFIG
20          The configuration file for the eigenface utility
21          instance
22      -m MODEL, --model MODEL
23          The model trained with this eigenface utility

1  usage: test.py [-h] [-i INPUT] [-m MODEL] [-c CONFIG] [-o OUTPUT]
2
3  We're assuming a <imageFileNameNoExt>.txt for eye position like
4  474 247 607 245
5  Comment line starting with # will be omitted
6  Or we'll use OpenCV's haarcascade detector to try locating eyes' positions
7
8  Note that if you want to save the recognition result
9  pass -o argument to specify the output file name
10 It's highly recommended to do so since OpenCV can't even draw large window
    properly ...
11
12 optional arguments:
13     -h, --help            show this help message and exit
14     -i INPUT, --input INPUT
15             The image we want to reconstruct and recognize on
16     -m MODEL, --model MODEL
17             The model trained with this eigenface utility
18     -c CONFIG, --config CONFIG
19             The configuration file for the eigenface utility
20             instance
21     -o OUTPUT, --output OUTPUT
22             The output file to save the reconstruction/recognition
23             result. If not specified, the program WILL NOT SAVE THE
24             RESULT

```

An example run would be:

```

1  python train.py -p "MyDataSet" -i .jpg -t .txt -c builtin.json -m model.color.npz
2
3  python test.py -i image_0318.jpg -m model.color.npz -c builtin.json -o
    similar.png

```

## Dataset

Our **EigenFaceUtils** can operate on the given dataset and also our own.

We constructed our dataset by labelling the eye position of *Caltec Database* and adding a few of the developer's face.

It was pretty awkward...

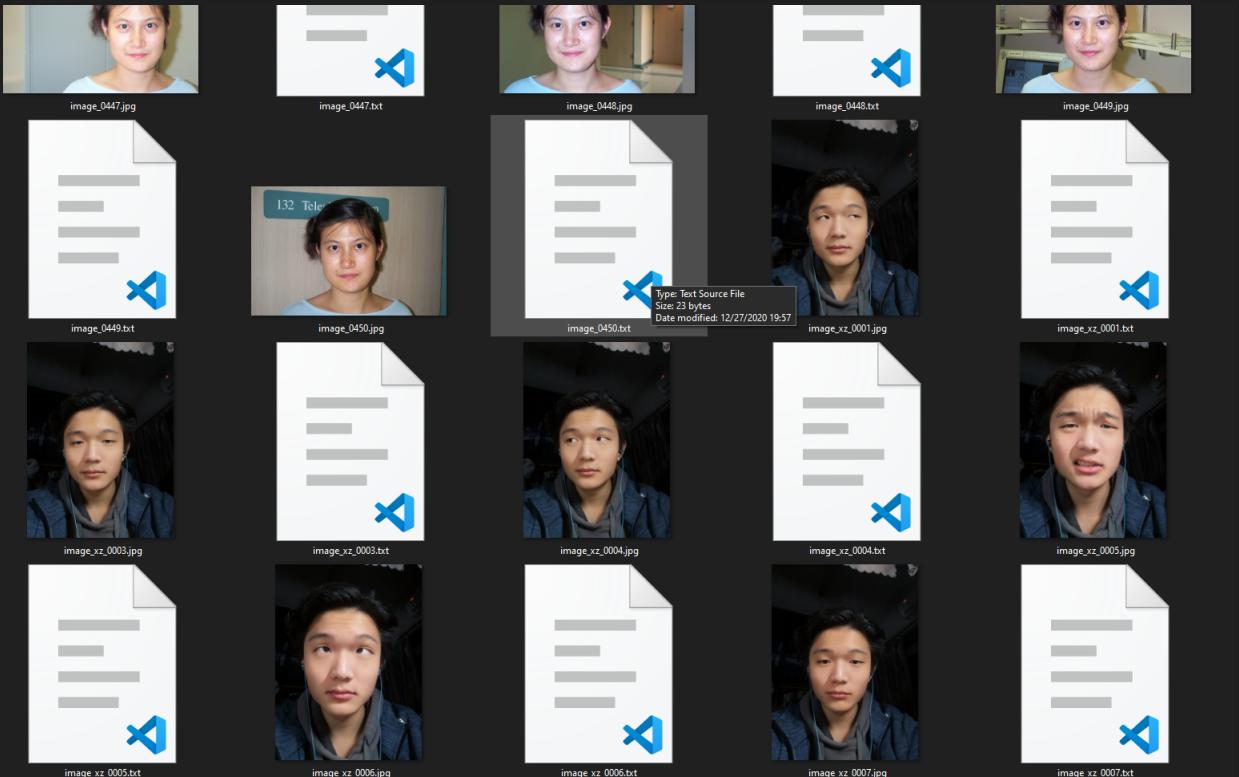
To make our life simplier, we only selected those images in *Caltec Database* that has a recognizable eyes to OpenCV's haar cascade recognizer.

And hand labelled the developer's eyes.

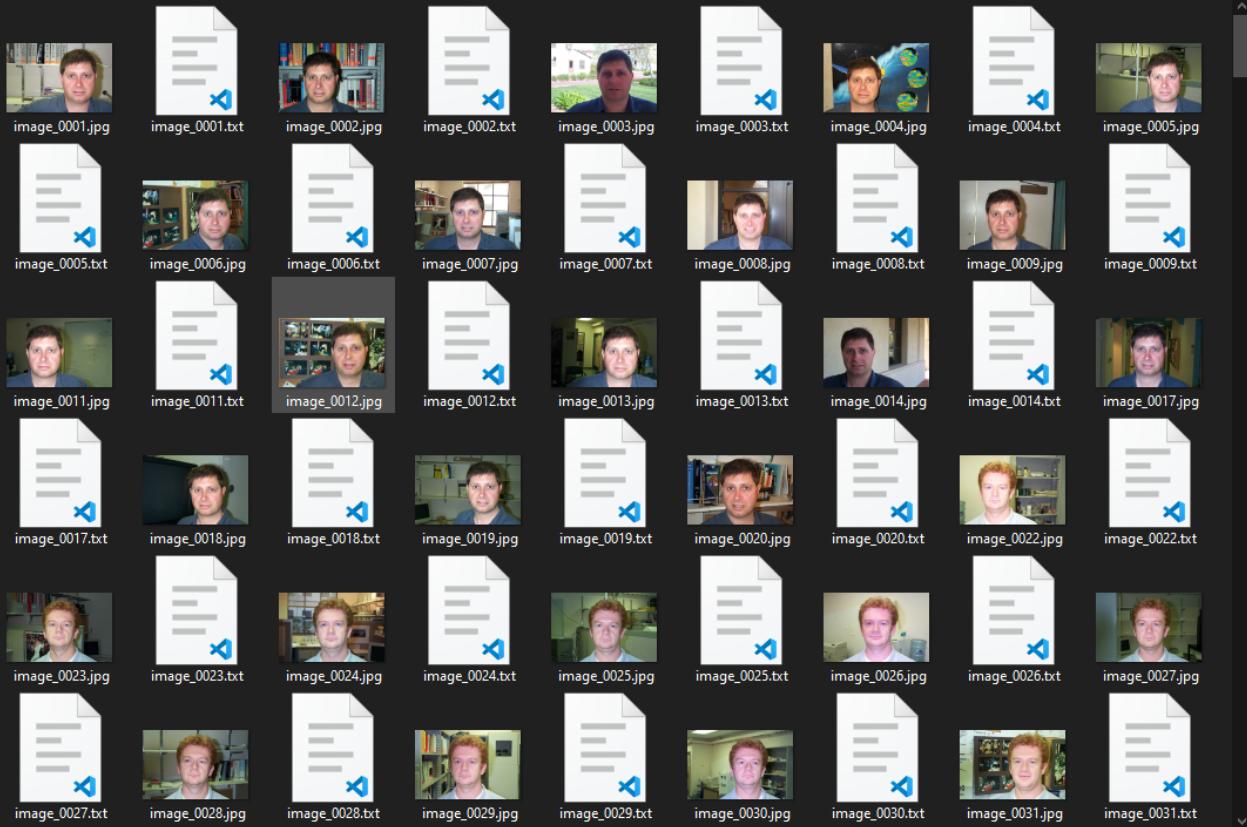
We provided 7 images of the developer, with label:

```
1 "image_xz_0001.jpg", "image_xz_0001.txt"
2 "image_xz_0002.jpg", "image_xz_0002.txt"
3 "image_xz_0003.jpg", "image_xz_0003.txt"
4 "image_xz_0004.jpg", "image_xz_0004.txt"
5 "image_xz_0005.jpg", "image_xz_0005.txt"
6 "image_xz_0006.jpg", "image_xz_0006.txt"
7 "image_xz_0007.jpg", "image_xz_0007.txt"
```

You can find the dataset in **MyDataSet** directory under the source directory



All other images in this directory are also labelled.



## Experiments And Results

### Training

We trained our eigenface model using this setup:

```
1  {
2      "width": 512,
3      "height": 512,
4      "left": [
5          188,
6          188
7      ],
8      "right": [
9          324,
10         188
11     ],
12      "isColor": true,
13      "nEigenFaces": 1000,
14      "targetPercentage": null,
15      "useBuiltIn": true,
16      "useHighgui": true
17  }
```

And we used this command on a Ubuntu machine:

```
1  ./train.py -p "MyDataSet" -i .jpg -t .txt -c builtin.json -m
model.MyDataSet512×512×3.npz
```

with hardware specification:

```
1          .+/o+-      xzdd@xzdd-ubuntu
2          yyyyy- -yyyyyy+  OS: Ubuntu 20.04 focal
3          ://+/////-yyyyyyo  Kernel: x86_64 Linux 5.4.0-58-generic
4          .++ .:/+++++/-.+sss/`  Uptime: 3d 13h 49m
5          .:+o: /+++++++/:-:/-  Packages: 1909
6          o:+o:+. ` .. ^`-/-oo++++/  Shell: zsh 5.8
7          .:+o:+o/.           `+sssooo+/  Disk: 338G / 670G (54%)
8          .++/+o+o: `           /sssooo.  CPU: Intel Core i7-9700 @ 8x 4.7GHz
[34.0°C]
9          /+++//+:`oo+o       /::--:.  GPU: GeForce RTX 2080 Ti
10         \+/+o+++'o++o       +++++.  RAM: 951MiB / 15921MiB
11         .++o++oo+:`        /dddhhh.
12         .+o+oo:.           `odhhhhh+
13         \+.++o+o`-^`.:ohdhhhhh+
14         `:o++`ohhhhhhhyo++os:
15         .o: `syhhhhhhh/.oo++o`  /
16         /osyyyyyyo++ooo++/  `-----+oo++o\:
17                           `oo++.
18
```

This is the program loading images and aligning them to our mask.

```
(py38) ~ homework3 git:(master) ./train.py -p "MyDataSet" -i .jpg -t .txt -c builtin.json -m model.MyDataSet512x512x3.npz
2020-12-28 11:03:51 xzdd-ubuntu eigenface[122502] INFO Loading configuration from builtin.json, with content: {'width': 512, 'height': 512, 'left': [188, 188], 'right': [324, 188], 'isColor': True, 'nEigenFaces': 1000, 'targetPercentage': None, 'useBuiltIn': True, 'useHighgui': True}
Processing batch: 345it [00:08, 8.14it/s]
```

This is the program computing eigenvectors/eigenvalues using **PCA**

```
(py38) ~ homework3 git:(master) ./train.py -p "MyDataSet" -i .jpg -t .txt -c builtin.json -m model.MyDataSet512x512x3.npz
2020-12-28 11:03:51 xzdd-ubuntu eigenface[122502] INFO Loading configuration from builtin.json, with content: {'width': 512, 'height': 512, 'left': [188, 188], 'right': [324, 188], 'isColor': True, 'nEigenFaces': 1000, 'targetPercentage': None, 'useBuiltIn': True, 'useHighgui': True}
Processing batch: 345it [01:06, 5.20it/s]
2020-12-28 11:04:58 xzdd-ubuntu eigenface[122502] INFO Getting 345 names and 345 batch
2020-12-28 11:04:58 xzdd-ubuntu eigenface[122502] INFO Beginning builtin PCACompute for 1000 eigenvalues/eigenvectors
```

Note that we set **targetPercentage** to **null** so we'll use **nEigenFaces: 1000** to try computing 1000 eigenfaces (possibly  $512 \times 512 \times 3 = 786432$ )

Eventually we can only get 345 eigenfaces:

```
(py38) ~ homework3 git:(master) ./train.py -p "MyDataSet" -i .jpg -t .txt -c builtin.json -m model.MyDataSet512x512x3.npz
2020-12-28 11:03:51 xzdd-ubuntu eigenface[122502] INFO Loading configuration from builtin.json, with content: {'width': 512, 'height': 512, 'left': [188, 188], 'right': [324, 188], 'isColor': True, 'nEigenFaces': 1000, 'targetPercentage': None, 'useBuiltIn': True, 'useHighgui': True}
Processing batch: 345it [01:06, 5.20it/s]
2020-12-28 11:04:58 xzdd-ubuntu eigenface[122502] INFO Getting 345 names and 345 batch
2020-12-28 11:04:58 xzdd-ubuntu eigenface[122502] INFO Beginning builtin PCACompute for 1000 eigenvalues/eigenvectors
2020-12-28 11:06:08 xzdd-ubuntu eigenface[122502] INFO Getting mean vectorized face: [[104.96231884 115.71594203 113.35072464 ... 123.13913043 117.17101449
106.61449275]] with shape: (1, 786432)
2020-12-28 11:06:08 xzdd-ubuntu eigenface[122502] INFO Getting sorted eigenvectors:
[[ 1.51637383e-03  1.5272011e-03  1.65011839e-03 ... -1.59191082e-06
  1.19530430e-04  2.64387095e-04]
 [-6.25459419e-04 -6.74333202e-04 -7.32235834e-04 ...  3.66228402e-03
  3.86894349e-03  3.52330871e-03]
 [-1.14571354e-03 -1.15501789e-03 -1.01845025e-03 ...  3.27172565e-04
  1.00252147e-04 -2.35212609e-05]
 ...
 [ 4.36742887e-04 -2.96260371e-04 -1.22413086e-03 ... -1.0280926e-03
 -8.18479782e-04 -9.07490186e-04]
 [ 1.93892789e-04  1.39758695e-04  1.53280597e-04 ... -1.36009244e-03
 -1.30934660e-03 -2.01625493e-03]
 [-1.76967576e-03 -9.32994634e-04 -6.18863852e-04 ... -1.55962110e-03
 -1.26607964e-03 -1.51802196e-03]]
of shape: (345, 786432)
FaceDict: 97%
```

And this is the program compressing image data using eigenfaces

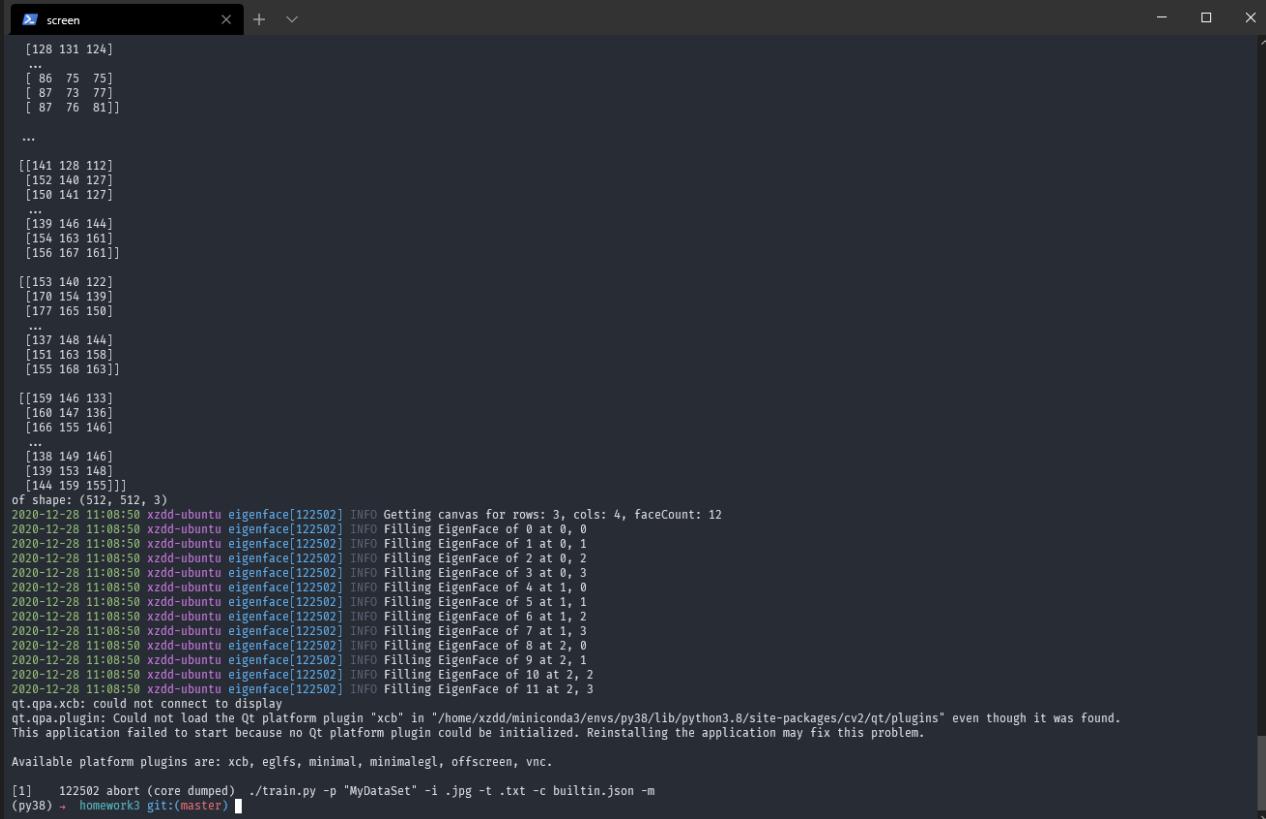
The compressed image data, from  $((512 \times 512 \times 3)^2 \times 8$  bits to  $345 \times 512 \times 512 \times 3 \times 64$  bit (or we could even use 8 bit values), gets a compression ratio of  $345 \times 8 / (512 \times 512 \times 3) = 0.0035$

Then it tries to save the model to the file we specified in the commandline arguments

```
(py38) + homework3 git:(master) ./train.py -p "MyDataSet" -i .jpg -t .txt -c builtin.json -m model.MyDataSet512x512x3.npz
2020-12-28 11:03:51 xzdd-ubuntu eigenface[122502] INFO Loading configuration from builtin.json, with content: {'width': 512, 'height': 512, 'left': [188, 188], 'right': [324, 188], 'isColor': True, 'nEigenFaces': 1000, 'targetPercentage': None, 'useBuiltin': True, 'useHighgui': True}
Processing batch: 345it [01:06, 5.20it/s]
2020-12-28 11:04:58 xzdd-ubuntu eigenface[122502] INFO Getting 345 names and 345 batch
2020-12-28 11:04:58 xzdd-ubuntu eigenface[122502] INFO Beginning builtin PCACompute for 1000 eigenvalues/eigenvectors
2020-12-28 11:06:08 xzdd-ubuntu eigenface[122502] INFO Getting mean vectorized face: [[104.96231884 115.71594203 113.35072464 ... 123.13913043 117.17101449 106.614492751]
...
2020-12-28 11:06:08 xzdd-ubuntu eigenface[122502] INFO Getting sorted eigenvectors:
[[ 1.51637383e-03  1.52720111e-03  1.65011839e-03 ... -1.59191082e-06
  1.19530430e-04  2.64387095e-04]
[-6.25459418e-04 -6.74333024e-04 -7.32235834e-04 ...  3.66228402e-03
  3.86894434e-03  3.52330871e-03]
[-1.14571135e-03 -1.15501789e-03 -1.01845025e-03 ...  3.27172565e-04
  1.00252147e-04 -2.35212609e-05]
...
[-4.36742887e-04 -2.96260371e-04 -1.22413086e-03 ... -1.02809826e-03
  -8.18479782e-04 -9.07450186e-04]
[ 1.93892789e-04  1.39758695e-04  1.53280597e-04 ... -1.36009244e-03
  -1.30934660e-03 -2.01625493e-03]
[-1.76967576e-03 -9.32994634e-04 -6.18863852e-04 ... -1.55962110e-03
  -1.26697964e-03 -1.51802196e-03]]
of shape: (345, 786432)
FaceDict: 100% | 345/345 [01:33<00:00, 3.68it/s]
2020-12-28 11:07:41 xzdd-ubuntu eigenface[122502] INFO Got face dict of length 345
2020-12-28 11:07:41 xzdd-ubuntu eigenface[122502] INFO Saving model: model.MyDataSet512x512x3.npz
```

And this is what happens when no GUI is booted up and we tried to show windows using

### cv2.imshow



Fortunately, we'll save the images to files before loading all of them to the screen.

So if your running on a server without a display, this is the **success** message for our program

Let's checkout the result:

1. We specified a target number of eigenfaces of 1000

But only 345 valid eigenvalues are computed, this is pretty reasonable

2. After alignment and color histogram equlization, we can get a mean face:



3. This is the first few eigenfaces:



The mean eigenface:



And we'd also generate a file with all the eigenfaces for you to exam:

But it's a little bit too large to be displayed here:

Property	Value
Origin	
Date taken	
Image	
Dimensions	16896 x 15872
Width	16896 pixels
Height	15872 pixels
Bit depth	8
File	
Name	alleigenfaces.png
Item type	PNG File
Folder path	C:\Users\56295\Desktop\CompVision...
Date created	12/22/2020 20:43
Date modified	12/22/2020 20:43
Size	143 MB
Attributes	A
Availability	
Offline status	
Shared with	

```
4. We've got a model of size: -rw-rw-r-- 1 xzdd xzdd 2.0G Dec 28 11:08  
model.MyDataSet512x512x3.npz
```

## Testing

Then, with the model and corresponding configuration file, we can use the eigenfaces we computed to compress our own image and get the reconstruction.

Then by comparing the Euclidean Distance between the projected test image with that of the ones in the database (saved to model), we can get the most similar face from the database and display (or render to file)

This is what we're going to do:

1. Load the model, including all eigenvectors(eigenfaces), the mean face vector and all compressed image data
2. Load the test image and go through the typical preprocessing procedure of alignment and histogram equilization
3. Do a matrix multiplication between the eigenvectors and the loaded (processed) test image to retrieve the weight (compressed version of the test image)
4. Do a Euclidean Distance comparison between all the images in the database (compressed using eigenface) to find the most similar image in the database (with smallest Euclidean Distance)
5. Reconstruct the loaded test image using eigenvectors:

```
1 img = mean  
2 img += np.mul(np.transpose(vectors), weights)
```

6. Reconstruct the most similar image in the database using the same method decribed above
7. Render those image to a canvas and label them, providing with necessary name attributes for you to view them
8. The rendered image are ordered like:

Original::Reconstructed Original::Most Weighted Eigenface::Reconstructed Most Similar Database Image::Original Most Similar Database Images

Testing with an unseen image (me):

```
(py38) → homework3 git:(master) ./test.py -i image_xz_0002.jpg -m model.MyDataSet512x512x3.npz -c builtin.json -o similarxz0002.png
2020-12-28 15:20:27 xzdd-ubuntu eigenface[147319] INFO Loading configuration from builtin.json, with content: {'width': 512, 'height': 512, 'left': [188, 188], 'right': [324, 188], 'isColor': True, 'nEigenFaces': 1000, 'targetPercentage': None, 'use_builtin': True, 'useHighgui': True}
2020-12-28 15:20:27 xzdd-ubuntu eigenface[147319] INFO Loading model: model.MyDataSet512x512x3.npz
2020-12-28 15:20:36 xzdd-ubuntu eigenface[147319] INFO Getting mean vectorized face: [[104.96231884 115.71594203 113.35072464 ... 123.13913043 117.17101449
106.61449275]] with shape: (1, 786432)
2020-12-28 15:20:36 xzdd-ubuntu eigenface[147319] INFO Getting sorted eigenvectors:
[[ 1.51637383e-03 1.52720111e-03 1.65011839e-03 ... -1.59191082e-06
1.19530430e-04 2.64387095e-04]
[-6.25459418e-04 -6.74333024e-04 -7.32235834e-04 ... 3.66228402e-03
3.86894434e-03 3.52330871e-03]
[-1.14571354e-03 -1.15501789e-03 -1.01845025e-03 ... 3.27172565e-04
1.00251247e-04 -2.35212609e-05]
...
[-4.36742887e-04 -2.96260371e-04 -1.22413086e-03 ... -1.02809826e-03
-8.18479782e-04 -9.07450186e-04]
[ 1.93892789e-04 1.39758695e-04 1.53280597e-04 ... -1.36009244e-03
-1.30934660e-03 -2.01625493e-03]
[-1.76967576e-03 -9.32994634e-04 -6.18863852e-04 ... -1.55962110e-03
-1.26697964e-03 -1.51802196e-03]]
of shape: (345, 786432)
2020-12-28 15:20:36 xzdd-ubuntu eigenface[147319] INFO Getting face dict of length: 345
2020-12-28 15:20:36 xzdd-ubuntu __main__[147319] INFO Found text file
2020-12-28 15:20:36 xzdd-ubuntu eigenface[147319] INFO Processing: image_xz_0002.txt
2020-12-28 15:20:36 xzdd-ubuntu __main__[147319] INFO Loading image: image_xz_0002.jpg
2020-12-28 15:20:36 xzdd-ubuntu eigenface[147319] INFO Getting eyes: [[1514. 2780.]
[2387. 2840.]]
2020-12-28 15:20:36 xzdd-ubuntu eigenface[147319] INFO Getting faceVect: [-873. -60.] and maskVect: [-136 0]
2020-12-28 15:20:36 xzdd-ubuntu eigenface[147319] INFO Getting faceNorm: 875.0594265534198 and maskNorm: 136.0
2020-12-28 15:20:36 xzdd-ubuntu eigenface[147319] INFO Should scale the image to: 0.155418016049105
2020-12-28 15:20:36 xzdd-ubuntu eigenface[147319] INFO Should rotate the image: -90.0 - -93.931671480022225 = 3.9316714800222456 degrees
2020-12-28 15:20:36 xzdd-ubuntu eigenface[147319] INFO Getting faceCenter: [1950.5 2810. ] and maskCenter: [256. 188.]
2020-12-28 15:20:36 xzdd-ubuntu eigenface[147319] INFO Should translate the image using: [-1694.5 -2622. ]
2020-12-28 15:20:36 xzdd-ubuntu eigenface[147319] INFO Getting # of channels: 3
2020-12-28 15:20:36 xzdd-ubuntu eigenface[147319] INFO Shape of eigenvectors and flat: (345, 786432), (786432, 1)
Recognizing: 100% |██████████| 345/345 [00:00<00:00, 95779.38it/s]
2020-12-28 15:20:37 xzdd-ubuntu eigenface[147319] INFO MOST SIMILAR FACE: MyDataSet/image_xz_0007.jpg WITH RESULT 22112.468775792233
2020-12-28 15:20:37 xzdd-ubuntu eigenface[147319] INFO Shape of flat: (786432, 1)
2020-12-28 15:20:37 xzdd-ubuntu eigenface[147319] INFO Processing: MyDataSet/image_xz_0007.txt
2020-12-28 15:20:37 xzdd-ubuntu eigenface[147319] INFO Getting eyes: [[1514. 2780.]
[2422. 2797.]]
2020-12-28 15:20:37 xzdd-ubuntu eigenface[147319] INFO Getting faceVect: [-901. -29.] and maskVect: [-136 0]
2020-12-28 15:20:37 xzdd-ubuntu eigenface[147319] INFO Getting faceNorm: 901.4665828526313 and maskNorm: 136.0
2020-12-28 15:20:37 xzdd-ubuntu eigenface[147319] INFO Should scale the image to: 0.1508652706455707
2020-12-28 15:20:37 xzdd-ubuntu eigenface[147319] INFO Should rotate the image: -90.0 - -91.8435118554498615 degrees
2020-12-28 15:20:37 xzdd-ubuntu eigenface[147319] INFO Getting faceCenter: [1971.5 2782.5] and maskCenter: [256. 188.]
2020-12-28 15:20:37 xzdd-ubuntu eigenface[147319] INFO Should translate the image using: [-1715.5 -2594.5]
2020-12-28 15:20:37 xzdd-ubuntu eigenface[147319] INFO Getting # of channels: 3
2020-12-28 15:20:37 xzdd-ubuntu eigenface[147319] INFO Successfully loaded the original image: MyDataSet/image_xz_0007.jpg
2020-12-28 15:20:37 xzdd-ubuntu __main__[147319] INFO Saving output to similarxz0002.png
qt.qpa.xcb: could not connect to display
qt.qpa.plugin: Could not load the Qt platform plugin "xcb" in "/home/xzdd/miniconda3/envs/py38/lib/python3.8/site-packages/cv2/qt/plugins" even though it was found.
This application failed to start because no Qt platform plugin could be initialized. Reinstalling the application may fix this problem.

Available platform plugins are: xcb, eglfs, minimal, minimalegl, offscreen, vnc.

[1] 147319 abort (core dumped) ./test.py -i image_xz_0002.jpg -m model.MyDataSet512x512x3.npz -c builtin.json -o similarxz0002.png
(py38) → homework3 git:(master) |
```

The command we used were:

```
1 ./test.py -i image_xz_0002.jpg -m model.MyDataSet512x512x3.npz -c builtin.json -o similarxz0002.png
```

You can see that we were using a server so no windows could be opened.

But luckily the images are saved to the directory, so you can check them via either some service on your server or **scp** them back to your local directory



We rendered the images at full size, so this should be a 512 by 5x512 image

You can see the reconstructed test image do resembles me in some way.

And it corrected my closed eyes (due to the fact that most of the images in the dataset are open-eyed) and tilted head (due to unaligned labelling)

Moreover, even my clothes are being reconstructed

On the right side is the most similar database image to the one being tested and the person on that is indeed me.

The most astonishing part is that: the reconstructed database image is almost identical to the original one, meaning the compression loses almost no data at all, even though the compression ratio is already 0.0035 (without the whole model)

During the testing, you can see the speed of recognizing the face is indeed pretty fast

(Even though, due to the fact that we saved our eigenvectors using 64-bit float values with compression, the model file itself is pretty large and chunky to load from disk)

Another test result:



Testing with an existing image:

```
(py38) → homework3 git:(master) ./test.py -i image_xz_0006.jpg -m model.MyDataSet512x512x3.npz -c builtin.json -o similarxz0006.png
2020-12-28 15:18:25 xzdd-ubuntu eigenface[146892] INFO Loading configuration from builtin.json, with content: {'width': 512, 'height': 512, 'left': [188, 188], 'right': [324, 188], 'isColor': True, 'nEigenFaces': 1000, 'targetPercentage': None, 'useBuiltin': True, 'useHighgui': True}
2020-12-28 15:18:25 xzdd-ubuntu eigenface[146892] INFO Loading model: model.MyDataSet512x512x3.npz
2020-12-28 15:18:34 xzdd-ubuntu eigenface[146892] INFO Getting mean vectorized face: [[104.96231884 115.71594203 113.35072464 ... 123.13913043 117.17101449
106.61449275]] with shape: (1, 786432)
2020-12-28 15:18:34 xzdd-ubuntu eigenface[146892] INFO Getting sorted eigenvectors:
[[ 1.51637383e-03 1.52720111e-03 1.65011839e-03 ... -1.59191082e-06
1.19530430e-04 2.64387095e-04]
[-6.25459418e-04 -6.74333024e-04 -7.32235834e-04 ... 3.66228402e-03
3.86894434e-03 3.52330871e-03]
[-1.14571354e-03 -1.15501789e-03 -1.01845025e-03 ... 3.27172565e-04
1.00252147e-04 -2.35212609e-05]
...
[-4.36742887e-04 -2.96260371e-04 -1.22413086e-03 ... -1.02809826e-03
-8.18479782e-04 -9.07450186e-04]
[ 1.93892789e-04 1.39758695e-04 1.53280597e-04 ... -1.36009244e-03
-1.30934660e-03 -2.01625493e-03]
[-1.76967576e-03 -9.32994634e-04 -6.18863852e-04 ... -1.55962110e-03
-1.26697964e-03 -1.51802196e-03]]
of shape: (345, 786432)
2020-12-28 15:18:34 xzdd-ubuntu eigenface[146892] INFO Getting face dict of length: 345
2020-12-28 15:18:34 xzdd-ubuntu __main__[146892] INFO Found text file
2020-12-28 15:18:34 xzdd-ubuntu eigenface[146892] INFO Processing: image_xz_0006.txt
2020-12-28 15:18:34 xzdd-ubuntu __main__[146892] INFO Loading image: image_xz_0006.jpg
2020-12-28 15:18:34 xzdd-ubuntu eigenface[146892] INFO Getting eyes: [[1364. 2439.]
[2551. 2497.]]
2020-12-28 15:18:34 xzdd-ubuntu eigenface[146892] INFO Getting faceVect: [-1187. -58.] and maskVect: [-136 0]
2020-12-28 15:18:34 xzdd-ubuntu eigenface[146892] INFO Getting faceNorm: 1188.4161728956738 and maskNorm: 136.0
2020-12-28 15:18:34 xzdd-ubuntu eigenface[146892] INFO Should scale the image to: 0.11443802524886952
2020-12-28 15:18:34 xzdd-ubuntu eigenface[146892] INFO Should rotate the image: -90.0 - -92.79740037767522 = 2.797400377675217 degree
s
2020-12-28 15:18:34 xzdd-ubuntu eigenface[146892] INFO Getting faceCenter: [1957.5 2468. ] and maskCenter: [256. 188.]
2020-12-28 15:18:34 xzdd-ubuntu eigenface[146892] INFO Should translate the image using: [-1701.5 -2280. ]
2020-12-28 15:18:34 xzdd-ubuntu eigenface[146892] INFO Getting # of channels: 3
2020-12-28 15:18:34 xzdd-ubuntu eigenface[146892] INFO Shape of eigenvectors and flat: (345, 786432), (786432, 1)
Recognizing: 100% |██████████| 345/345 [00:00<00:00, 93098.81it/s]
2020-12-28 15:18:34 xzdd-ubuntu eigenface[146892] INFO MOST SIMILAR FACE: MyDataSet/image_xz_0006.jpg WITH RESULT 0.0
2020-12-28 15:18:35 xzdd-ubuntu eigenface[146892] INFO Shape of flat: (786432, 1)
2020-12-28 15:18:35 xzdd-ubuntu eigenface[146892] INFO Processing: MyDataSet/image_xz_0006.txt
2020-12-28 15:18:35 xzdd-ubuntu eigenface[146892] INFO Getting eyes: [[1364. 2439.]
[2551. 2497.]]
2020-12-28 15:18:35 xzdd-ubuntu eigenface[146892] INFO Getting faceVect: [-1187. -58.] and maskVect: [-136 0]
2020-12-28 15:18:35 xzdd-ubuntu eigenface[146892] INFO Getting faceNorm: 1188.4161728956738 and maskNorm: 136.0
2020-12-28 15:18:35 xzdd-ubuntu eigenface[146892] INFO Should scale the image to: 0.11443802524886952
2020-12-28 15:18:35 xzdd-ubuntu eigenface[146892] INFO Should rotate the image: -90.0 - -92.79740037767522 = 2.797400377675217 degree
s
2020-12-28 15:18:35 xzdd-ubuntu eigenface[146892] INFO Getting faceCenter: [1957.5 2468. ] and maskCenter: [256. 188.]
2020-12-28 15:18:35 xzdd-ubuntu eigenface[146892] INFO Should translate the image using: [-1701.5 -2280. ]
2020-12-28 15:18:35 xzdd-ubuntu eigenface[146892] INFO Getting # of channels: 3
2020-12-28 15:18:35 xzdd-ubuntu eigenface[146892] INFO Successfully loaded the original image: MyDataSet/image_xz_0006.jpg
2020-12-28 15:18:35 xzdd-ubuntu __main__[146892] INFO Saving output to similarxz0006.png
qt.qpa.xcb: could not connect to display
qt.qpa.plugin: Could not load the Qt platform plugin "xcb" in "/home/xzdd/miniconda3/envs/py38/lib/python3.8/site-packages/cv2/qt/plugins" even though it was found.
This application failed to start because no Qt platform plugin could be initialized. Reinstalling the application may fix this problem.

Available platform plugins are: xcb, eglfs, minimal, minimalegl, offscreen, vnc.

[1] 146892 abort (core dumped) ./test.py -i image_xz_0006.jpg -m model.MyDataSet512x512x3.npz -c builtin.json
```

## loading model

### image already in database

From the existing images' result, we can also confirm that the compression is indeed enough for accuracy and reducing dimension can significantly reduce the time required to do the computation of other kinds (utilizing change of basis)



## Additional Experiments

### Additional Training

Training with `isColor=false`, `targetPercentage=null` and `useBuiltIn=false`:

The configuration we used:

```
1  {
2      "width": 128,
3      "height": 128,
4      "left": [
5          48,
6          48
7      ],
8      "right": [
9          80,
10         48
11     ],
12     "isColor": false,
13     "nEigenFaces": 50,
14     "targetPercentage": null,
15     "useBuiltIn": false,
16     "useHighgui": true
17 }
```

Here we're using grayscale training and our own implementation of PCA

We also constructed a small dataset to speed things up: `MySmallDataSet`

Everyone has only three faces here.

Here the `targetPercentage` is set to null, so we won't have to compute all eigenvalues to try reaching the target information retain rate

```
(local38) → homework3 git:(master) ✘ python ./train.py -p "MySmallDataSet" -i .jpg -t .txt -c default.json -m model.MySm
allDataSet128x128x1.npz
2020-12-28 16:13:17 DESKTOP-XZDDSPC eigenface[22492] INFO Loading configuration from default.json, with content: {'width'
: 128, 'height': 128, 'left': [48, 48], 'right': [80, 48], 'isColor': False, 'nEigenFaces': 32, 'targetPercentage': None,
'useBuiltIn': False, 'useHighgui': True}
Processing batch: 45it [00:00, 59.37it/s]
2020-12-28 16:13:18 DESKTOP-XZDDSPC eigenface[22492] INFO Getting 45 names and 45 batch
2020-12-28 16:13:18 DESKTOP-XZDDSPC eigenface[22492] INFO Getting mean vectorized face: [[105.2           106.           104.7
333333 ... 115.8          116.71111111
115.44444444]] with shape: (1, 16384)
2020-12-28 16:13:18 DESKTOP-XZDDSPC eigenface[22492] INFO Trying to compute the covariance matrix
2020-12-28 16:13:22 DESKTOP-XZDDSPC eigenface[22492] INFO Getting covar of shape: (16384, 16384)
2020-12-28 16:13:22 DESKTOP-XZDDSPC eigenface[22492] INFO Getting covariance matrix:
[[ 5559.84545455  5422.97727273  4975.94090909 ... -1017.50454545
-1084.89545455 -1220.18181818]
[ 5422.97727273  5483.54545455  4942.95454545 ... -995.36363636
-1061.95454545 -1221.18181818]
[ 4975.94090909  4942.95454545  5227.38181818 ... -405.6
-709.96515152 -997.74242424]
...
[-1017.50454545 -995.36363636 -405.6       ... 8662.39090909
8570.55454545 8123.86363636]
[-1084.89545455 -1061.95454545 -709.96515152 ... 8570.55454545
8832.48282828 8476.54040404]
[-1220.18181818 -1221.18181818 -997.74242424 ... 8123.86363636
8476.54040404 8362.11616162]]
8 seconds to get 32 eigenvalues
vectors
2020-12-28 16:13:22 DESKTOP-XZDDSPC eigenface[22492] INFO Begin computing 32 eigenvalues/eigenvectors
2020-12-28 16:13:30 DESKTOP-XZDDSPC eigenface[22492] INFO Getting 32 eigenvalues and eigenvectors with shape (16384, 32)
2020-12-28 16:13:30 DESKTOP-XZDDSPC eigenface[22492] INFO Getting sorted eigenvalues:
[13692077.82648995 8065424.82527987 4945670.55159279 3883130.78841176
3030309.58099939 2835500.38371136 2427842.1930707 1911457.71950977
1779447.25859033 1573182.62766051 1456963.11799481 1359633.86157442
1155014.95786098 1056813.40248439 951043.12834829 872667.24346596
860681.71336261 767367.56220602 743076.28065638 666788.53745266
650917.3351369 594233.2315434 549146.13607777 522088.96079639
503104.61461041 475584.82595126 466131.64726078 438166.47782317
413767.387414 380024.56284485 360908.19535676 333803.59797119]
of shape: (32,)
2020-12-28 16:13:30 DESKTOP-XZDDSPC eigenface[22492] INFO Getting sorted eigenvectors:
[[ -1.22082472e-02 -1.22170198e-02 -1.15940454e-02 ... 3.55293140e-03
4.01347596e-03 5.11466732e-03]
[ -4.20471112e-03 -2.8452086e-03 -2.12906552e-03 ... 2.21606187e-02
2.15921188e-02 1.99257795e-02]
[ 1.20820334e-02 1.39370645e-02 9.93585603e-03 ... -2.05471490e-03
-2.12010685e-03 -2.91506808e-04]
...
[-2.58951293e-03 6.66316446e-04 3.51098408e-03 ... 2.07559640e-03
1.43094726e-03 -3.75370247e-05]
[ 3.16011767e-03 1.78737093e-03 -1.93208445e-03 ... -1.01141184e-02
-7.52251863e-03 -3.59788604e-03]
[-2.65106330e-02 -2.55080570e-02 -1.91679145e-02 ... -9.51263106e-03
-5.91214268e-03 -4.91148062e-03]]
of shape: (32, 16384)
FaceDict: 100%|| 45/45 [00:00<00:00, 4101.70it/s]
2020-12-28 16:13:31 DESKTOP-XZDDSPC eigenface[22492] INFO Got face dict of length 45
2020-12-28 16:13:31 DESKTOP-XZDDSPC eigenface[22492] INFO Saving model: model.MySmallDataSet128x128x1.npz
2020-12-28 16:13:31 DESKTOP-XZDDSPC eigenface[22492] INFO Model: model.MySmallDataSet128x128x1.npz saved
2020-12-28 16:13:31 DESKTOP-XZDDSPC eigenface[22492] INFO Computing eigenfaces
2020-12-28 16:13:31 DESKTOP-XZDDSPC eigenface[22492] INFO Getting eigenfaces of shape (32, 128, 128)
2020-12-28 16:13:31 DESKTOP-XZDDSPC eigenface[22492] INFO Getting mean eigenface
[[112 117 105 ... 111 111 89]
[ 83 91 94 ... 113 107 114]
[ 89 101 111 ... 129 123 112]
...
[158 131 99 ... 166 156 140]
[146 115 104 ... 168 180 163]
```

Note that computing **32 eigenvalues/eigenvectors** took **8 seconds**

Of course we'd get the basic results:

Mean Eigenfaces:



First 12 Eigenfaces:



Then we performed another training using `targetPercentage=0.95` to retain 0.95 of all the information

Training with `targetPercentage=0.95`

```
1  {
2      "width": 128,
3      "height": 128,
4      "left": [
5          48,
6          48
7      ],
8      "right": [
9          80,
10         48
11     ],
12      "isColor": false,
13      "nEigenFaces": 50,
14      "targetPercentage": 0.95,
15      "useBuiltIn": false,
16      "useHighgui": true
17 }
```

Other procedures are purely identical

```
(local38) ➔ homework3 git:(master) ✘ python ./train.py -p "MySmallDataSet" -i .jpg -t .txt -c default.json -m model.M
ySmallDataSet128x128x1.npz
2020-12-28 15:55:45 DESKTOP-XZDDSPC eigenface[14372] INFO Loading configuration from default.json, with content: {'width': 128, 'height': 128, 'left': [48, 48], 'right': [80, 48], 'isColor': False, 'nEigenFaces': 50, 'targetPercentage': 0.95, 'useBuiltIn': False, 'useHighgui': True}
Processing batch: 45it [00:00, 59.86it/s]
2020-12-28 15:55:46 DESKTOP-XZDDSPC eigenface[14372] INFO Getting 45 names and 45 batch
2020-12-28 15:55:46 DESKTOP-XZDDSPC eigenface[14372] INFO Getting mean vectorized face: [[105.2           106.           10
4.7333333 ... 115.8          116.7111111
115.44444444]] with shape: (1, 16384)
2020-12-28 15:55:46 DESKTOP-XZDDSPC eigenface[14372] INFO Trying to compute the covariance matrix
2020-12-28 15:55:49 DESKTOP-XZDDSPC eigenface[14372] INFO Getting covar of shape: (16384, 16384)
2020-12-28 15:55:49 DESKTOP-XZDDSPC eigenface[14372] INFO Getting covariance matrix:
[[ 5559.84545455  5422.97727273  4975.94090909 ... -1017.50454545
-1084.89545455 -1220.18181818]
[ 5422.97727273  5483.54545455  4942.95454545 ... -995.36363636
-1061.95454545 -1221.18181818]
[ 4975.94090909  4942.95454545  5227.38181818 ... -405.6
-709.96515152 -997.74242424]
...
[-1017.50454545 -995.36363636 -405.6       ... 8662.39090909
8570.55454545  8123.86363636]
[-1084.89545455 -1061.95454545 -709.96515152 ... 8570.55454545
8832.48282828  8470.54040404]
[-1220.18181818 -1221.18181818 -997.74242424 ... 8123.86363636
8476.54040404  8362.11616162]]
2020-12-28 15:55:49 DESKTOP-XZDDSPC eigenface[14372] INFO Begin computing all eigenvalues
2020-12-28 16:04:00 DESKTOP-XZDDSPC eigenface[14372] INFO Getting all eigenvalues:
[ 1.36920778e+07  8.06542483e+06  4.94567055e+06 ... -3.13223211e-09
-4.84075366e-09 -7.37173973e-09]
of shape: (16384,)
2020-12-28 16:04:00 DESKTOP-XZDDSPC eigenface[14372] INFO For a energy percentage of 0.95, we need 32 vectors from 16384
2020-12-28 16:04:00 DESKTOP-XZDDSPC eigenface[14372] INFO Begin computing 32 eigenvalues/eigenvectors
2020-12-28 16:04:09 DESKTOP-XZDDSPC eigenface[14372] INFO Getting 32 eigenvalues and eigenvectors with shape (16384, 32)
2020-12-28 16:04:09 DESKTOP-XZDDSPC eigenface[14372] INFO Getting sorted eigenvalues:
[13692077.82648995 8065424.82527987 4945670.55159279 3883130.78841176
3030309.58099939 2835500.38371136 2427842.1930707 1911457.71950977
1779447.25859033 1573182.62766051 1456963.11799481 1359633.86157442
1155014.95786098 1056813.40248439 951043.12834829 872667.24346596
860681.71336261 767367.56220602 743076.28065638 666788.53745266
650917.3351369 594323.2315434 549146.13607777 522088.96079639
503104.61461041 475584.82595126 466131.64726078 438166.47782317
413767.387414 380024.56284485 360908.19535676 333803.59797119]
of shape: (32,)
2020-12-28 16:04:09 DESKTOP-XZDDSPC eigenface[14372] INFO Getting sorted eigenvectors:
[[-1.22082472e-02 -1.22170198e-02 -1.15940454e-02 ... 3.55293140e-03
4.01347596e-03 5.11466732e-03]
[-4.20471112e-03 -2.84525086e-03 -2.12906552e-03 ... 2.21606187e-02
2.15921188e-02 1.99257795e-02]
[ 1.20820334e-02 1.39370645e-02 9.93585603e-03 ... -2.05471490e-03
-2.12010685e-03 -2.91506808e-04]
...
[-2.58951293e-03 6.66316446e-04 3.51098408e-03 ... 2.07559640e-03
1.43094726e-03 -3.75370247e-05]
[ 3.16011767e-03 1.78737093e-03 -1.93208445e-03 ... -1.01141184e-02
-7.52251863e-03 -3.59788604e-03]
[-2.65106330e-02 -2.55080570e-02 -1.91679145e-02 ... -9.51263106e-03
-5.91214268e-03 -4.91148062e-03]]
of shape: (32, 16384)
FaceDict: 100% | 45/45 [00:00<00:00, 5636.50it/s]
2020-12-28 16:04:09 DESKTOP-XZDDSPC eigenface[14372] INFO Got face dict of length 45
2020-12-28 16:04:09 DESKTOP-XZDDSPC eigenface[14372] INFO Saving model: model.MySmallDataSet128x128x1.npz
2020-12-28 16:04:09 DESKTOP-XZDDSPC eigenface[14372] INFO Model: model.MySmallDataSet128x128x1.npz saved
2020-12-28 16:04:09 DESKTOP-XZDDSPC eigenface[14372] INFO Computing eigenfaces
2020-12-28 16:04:09 DESKTOP-XZDDSPC eigenface[14372] INFO Getting eigenfaces of shape (32, 128, 128)
2020-12-28 16:04:09 DESKTOP-XZDDSPC eigenface[14372] INFO Getting mean eigenface
```

took 8 minutes to get all eigenvalues

And this means to compute that 0.95 ratio without providing the function with the actual number of eigenvectors we want, we'd wait an additional **8 minutes** for this small dataset with only 45 images and a processed size of 128x128x1

It's easy to notice that the results are the same



eigenmean.png



eigenmean32.png



eigenfaces.png



eigenfaces32.png

**Training with target eigenface number = 50**

```
(local38) → homework3 git:(master) ✘ python ./train.py -p "MySmallDataSet" -i .jpg -t .txt -c default.json -m model.MySmallDataSet128x128x1x50.npz
2020-12-28 16:26:38 DESKTOP-XZDDSPC eigenface[20856] INFO Loading configuration from default.json, with content: {'width': 128, 'height': 128, 'left': [48, 48], 'right': [80, 48], 'isColor': False, 'nEigenFaces': 50, 'targetPercentage': None, 'useBuiltIn': False, 'useHighgui': True}
Processing batch: 45it [00:00, 59.44it/s]
2020-12-28 16:26:39 DESKTOP-XZDDSPC eigenface[20856] INFO Getting 45 names and 45 batch
2020-12-28 16:26:39 DESKTOP-XZDDSPC eigenface[20856] INFO Getting mean vectorized face: [[105.2           106.          104.7
333333 ... 115.8           116.71111111
115.44444444]] with shape: (1, 16384)
2020-12-28 16:26:39 DESKTOP-XZDDSPC eigenface[20856] INFO Trying to compute the covariance matrix
2020-12-28 16:26:42 DESKTOP-XZDDSPC eigenface[20856] INFO Getting covar of shape: (16384, 16384)
2020-12-28 16:26:42 DESKTOP-XZDDSPC eigenface[20856] INFO Getting covariance matrix:
[[ 5559.84545455  5422.97727273  4975.94090909 ... -1017.50454545
-1084.89545455 -1220.18181818]
[ 5422.97727273  5483.54545455  4942.95454545 ... -995.36363636
-1061.95454545 -1221.18181818]
[ 4975.94090909  4942.95454545  5227.38181818 ... -405.6
-709.96515152 -997.74242424]
...
[-1017.50454545 -995.36363636 -405.6       ...  8662.39090909
8570.55454545  8123.86363636]
[-1084.89545455 -1061.95454545 -709.96515152 ...  8570.55454545
8832.48282828  8476.54040404]
[-1220.18181818 -1221.18181818 -97.74242424 ...  8476.86068068
8476.54040404  8362.11616162]]
2020-12-28 16:26:42 DESKTOP-XZDDSPC eigenface[20856] INFO Begin computing 50 eigenvalues/eigenvectors
2020-12-28 16:27:12 DESKTOP-XZDDSPC eigenface[20856] INFO Getting 50 eigenvalues and eigenvectors with shape (16384, 50)
2020-12-28 16:27:12 DESKTOP-XZDDSPC eigenface[20856] INFO Getting sorted eigenvalues:
[[ 1.36920778e+07  8.06542483e+06  4.94567055e+06  3.88313079e+06
3.03030958e+06  2.83550038e+06  2.42784219e+06  1.91145772e+06
1.77944726e+06  1.57318263e+06  1.45696312e+06  1.35963386e+06
1.15501496e+06  1.05681340e+06  9.51043128e+05  8.72667243e+05
8.60681713e+05  7.67367562e+05  7.43076281e+05  6.66788537e+05
6.50917335e+05  5.94323232e+05  5.49146136e+05  5.22088961e+05
5.03104615e+05  4.75584826e+05  4.66131647e+05  4.38166478e+05
4.13767387e+05  3.80024563e+05  3.60908195e+05  3.33803598e+05
3.21100430e+05  3.10183150e+05  2.98025904e+05  2.74127350e+05
2.48076632e+05  2.42620441e+05  2.30272108e+05  2.19841886e+05
2.03463411e+05  1.96540753e+05  1.82455940e+05  1.71103940e+05
5.92915720e-10  5.84871586e-10  5.50996692e-10  -5.57130828e-10
-5.59378057e-10 -5.86506010e-10]
of shape: (50,)
2020-12-28 16:27:12 DESKTOP-XZDDSPC eigenface[20856] INFO Getting sorted eigenvectors:
[[ 0.01220825  0.01221702  0.01159405 ... -0.00355293 -0.00401348
-0.00511467]
[ 0.00420471  0.00284525  0.00212907 ... -0.02216062 -0.02159212
-0.01992578]
[-0.01208203 -0.01393706 -0.00993586 ...  0.00205471  0.00212011
 0.00029151]
...
[-0.00622595 -0.00240548 -0.00289496 ...  0.00330015  0.00455527
 0.00284277]
[ 0.00318609  0.00860076 -0.00335022 ...  0.00025306 -0.003702
 0.01103413]
[ 0.03488897 -0.00829044  0.00933665 ...  0.00791596  0.00914027
-0.01174183]]
of shape: (50, 16384)
FaceDict: 100% |██████████| 45/45 [00:00<00:00, 5012.58it/s]
2020-12-28 16:27:12 DESKTOP-XZDDSPC eigenface[20856] INFO Got face dict of length 45
```

Training with target eigenface number = 100

```
(local38) → homework3 git:(master) X python ./train.py -p "MySmallDataSet" -i .jpg -t .txt -c default.json -m model.MySmal
allDataSet128x128x1x100.npz
2020-12-28 16:31:36 DESKTOP-XZDDSPC eigenface[23712] INFO Loading configuration from default.json, with content: {'width': 128, 'height': 128, 'left': [48, 48], 'right': [80, 48], 'isColor': False, 'nEigenFaces': 100, 'targetPercentage': None, 'useBuiltIn': False, 'useHighgui': True}
Processing batch: 45it [00:00, 59.64it/s]
2020-12-28 16:31:37 DESKTOP-XZDDSPC eigenface[23712] INFO Getting 45 names and 45 batch
2020-12-28 16:31:37 DESKTOP-XZDDSPC eigenface[23712] INFO Getting mean vectorized face: [[105.2           106.           104.7
333333 ... 115.8          116.71111111
 115.44444444]] with shape: (1, 16384)
2020-12-28 16:31:37 DESKTOP-XZDDSPC eigenface[23712] INFO Trying to compute the covariance matrix
2020-12-28 16:31:40 DESKTOP-XZDDSPC eigenface[23712] INFO Getting covar of shape: (16384, 16384)
2020-12-28 16:31:40 DESKTOP-XZDDSPC eigenface[23712] INFO Getting covariance matrix:
[[ 5559.84545455  5422.97727273  4975.94090909 ... -1017.50454545
 -1084.89545455 -1220.18181818]
 [ 5422.97727273  5483.54545455  4942.95454545 ... -995.36363636
 -1061.95454545 -1221.18181818]
 [ 4975.94090909  4942.95454545  5227.38181818 ... -405.6
 -709.96515152 -997.74242424]
 ...
 [-1017.50454545 -995.36363636 -405.6       ...  8662.39090909
 8570.55454545  8123.86363636]
 [-1084.89545455 -1061.95454545 -709.96515152 ...  8570.55454545
 8832.48282828  8476.54040404]
 [-1220.18181818 -1221.18181818 -1221.18181818 ...  8123.86363636
 8476.54040404  8362.11616162]]
2020-12-28 16:31:40 DESKTOP-XZDDSPC eigenface[23712] INFO Begin computing 100 eigenvalues/eigenvectors
2020-12-28 16:34:15 DESKTOP-XZDDSPC eigenface[23712] INFO Getting 100 eigenvalues and eigenvectors with shape (16384, 100)
2020-12-28 16:34:15 DESKTOP-XZDDSPC eigenface[23712] INFO Getting sorted eigenvalues:
[[ 1.36920778e+07  8.06542483e+06  4.94567055e+06  3.88313079e+06
 3.03030958e+06  2.83550038e+06  2.42784219e+06  1.91145772e+06
 1.77944726e+06  1.57318263e+06  1.45696312e+06  1.35963386e+06
 1.15501496e+06  1.05681340e+06  9.51043128e+05  8.72667243e+05
 8.60681713e+05  7.67367562e+05  7.43076281e+05  6.66788537e+05
 6.50917335e+05  5.94323232e+05  5.49146136e+05  5.22088961e+05
 5.03104615e+05  4.75584826e+05  4.66131647e+05  4.38166478e+05
 4.13767387e+05  3.80024563e+05  3.60908195e+05  3.33803598e+05
 3.21100430e+05  3.10183150e+05  2.98025904e+05  2.74127350e+05
 2.48076632e+05  2.42620441e+05  2.30272108e+05  2.19841886e+05
 2.03463411e+05  1.96540753e+05  1.82455940e+05  1.71103940e+05
 6.31881861e-10  6.10329092e-10  5.84871586e-10  5.83418051e-10
 5.82367694e-10  5.74516477e-10  5.64402973e-10  5.58294802e-10
 5.53794041e-10  5.52210686e-10  5.47673919e-10  5.42134368e-10
 5.39260513e-10  5.37138540e-10  5.34487323e-10  5.32825193e-10
 5.31187904e-10  5.29370739e-10  5.29056431e-10  5.25258509e-10
 5.18375149e-10  5.10876335e-10  5.10628239e-10  5.10518172e-10
 5.09951279e-10  5.06413436e-10  5.05659206e-10  5.05565110e-10
 -5.04237028e-10  -5.04512979e-10  -5.04940374e-10  -5.05032066e-10
 -5.10600410e-10  -5.10711723e-10  -5.11709239e-10  -5.14101278e-10
 -5.15957769e-10  -5.27134817e-10  -5.27475672e-10  -5.29417688e-10
 -5.29775130e-10  -5.30181806e-10  -5.36656609e-10  -5.37412488e-10
 -5.38496639e-10  -5.40407823e-10  -5.47553819e-10  -5.49078537e-10
 -5.53804308e-10  -5.54943987e-10  -5.77271024e-10  -5.78746313e-10
 -5.86506010e-10  -5.93412114e-10  -6.08456431e-10  -6.26581087e-10]
of shape: (100,)
2020-12-28 16:34:15 DESKTOP-XZDDSPC eigenface[23712] INFO Getting sorted eigenvectors:
[[[-0.01220825 -0.01221702 -0.01159405 ...  0.00355293  0.00401348
 0.00511467]
 [-0.00420471 -0.00284525 -0.00212907 ...  0.02216062  0.02159212
 0.01992578]
 [ 0.01208203  0.01393706  0.00993586 ... -0.00205471 -0.00212011
 -0.00029151]
 ...]
```

You'll notice even getting more usable eigenvectors is faster than computing all of the eigenvalues

## **Additional Testing**

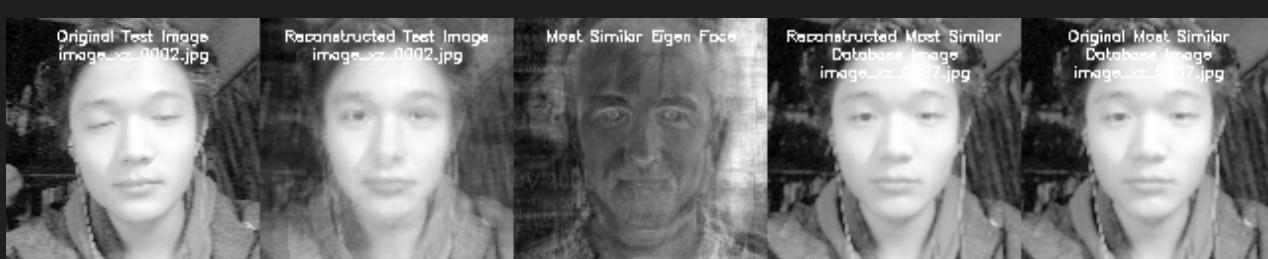
Then we performed the test described in the colored image section

32:





50:



100:



