

Homework 5

Homework 5

Basic Information

Implementation

Results

Training

Evaluation

LeNet-5

Dense

Prediction

Basic Information

- Student ID: 3180105504
- Student Name: Xu Zhen
- Instructor Name: Song Mingli
- Course Name: Computer Vision
- Homework Name: **CNN**

Implementation

We used **keras** to simplify the implementation (although we used **tensorflow.keras** eventually)

We used **tensorflow 2** as the **keras** backbone

And we used **matplotlib** to provide support for plotting and drawing

We also used some other small packages

```
1  #!/ python
2  import numpy as np
3  from scipy import interpolate
4  import matplotlib.pyplot as plt
5  import matplotlib as mpl
6
7  import tensorflow as tf
8  from tensorflow.keras.callbacks import TensorBoard, ModelCheckpoint, Callback,
    LearningRateScheduler, ReduceLROnPlateau
```

```

9  from tensorflow.keras.layers import Input, Dense, Flatten, Dropout, Activation,
   Conv2D, MaxPool2D, AveragePooling2D, BatchNormalization
10 from tensorflow.keras.models import Model, load_model, save_model
11 from tensorflow.keras.optimizers import Adam
12 from tensorflow.keras.losses import SparseCategoricalCrossentropy
13 from tensorflow.keras.metrics import SparseCategoricalAccuracy
14 import tensorflow_datasets as tfds
15
16 import coloredlogs
17 import logging

```

`matplotlib` and `coloredlogs` need to be configured:

```

1  # initialization for plotting and logging
2  # Setting up font for matplotlib
3  mpl.rc("font", family=["Josefin Sans", "Consolas", "Ubuntu", "Fira Code",
   "Inconsolata"], weight="medium", style="italic")
4  plt.style.use('dark_background')
5
6  coloredlogs.install("INFO")
7  log = logging.getLogger(__name__)

```

To prevent `tensorflow` from consuming all of our GPU memory (in practice it needs only like 900MB of GPU Memory and we have a 2080ti with 11GB Memory), we tell it to grow its memory usage:

```

1  def growth():
2      gpus = tf.config.experimental.list_physical_devices('GPU')
3      if gpus:
4          try:
5              # Currently, memory growth needs to be the same across GPUs
6              for gpu in gpus:
7                  tf.config.experimental.set_memory_growth(gpu, True)
8              logical_gpus = tf.config.experimental.list_logical_devices('GPU')
9              log.info(f"{len(gpus)} Physical GPUs, {len(logical_gpus)} Logical
   GPUs")
10         except RuntimeError as e:
11             # Memory growth must be set before GPUs have been initialized
12             log.error(e)

```

After calling this function, you can see that the GPU memory will NOT be fully consumed:

```
(py38) → homework5 git:(master) X nvidia-smi
Sat Jan 16 13:46:21 2021
```

```
(py38) → homework5 git:(master) X
```

1. Classical LeNet-5 with convolutional layers

```
1 def lenet5(input_shape, n_classes, lr):
2
3     inputs = Input(shape=input_shape)
4     model = Conv2D(6, kernel_size=(5, 5), strides=(1, 1),
5     activation='tanh', input_shape=input_shape, padding="same")(inputs)
6     model = AveragePooling2D(pool_size=(2, 2), strides=(2, 2),
7     padding='valid')(model)
```

```

6     model = Conv2D(16, kernel_size=(5, 5), strides=(1, 1),
activation='tanh', padding='valid')(model)
7     model = AveragePooling2D(pool_size=(2, 2), strides=(2, 2),
padding='valid')(model)
8     model = Flatten()(model)
9     model = Dense(120, activation='tanh')(model)
10    model = Dense(84, activation='tanh')(model)
11    model = Dense(n_classes, activation='softmax')(model)
12    model = Model(inputs, model)
13    model.compile(
14        optimizer=Adam(lr=lr),
15        loss=SparseCategoricalCrossentropy(from_logits=True),
16        metrics=[SparseCategoricalAccuracy()])
17    return model

```

Model summary:

1	Model: "model"		
2			
3	Layer (type)	Output Shape	Param #
4			
5	input_1 (InputLayer)	[(None, 28, 28, 1)]	0
6			
7	conv2d (Conv2D)	(None, 28, 28, 6)	156
8			
9	average_pooling2d (AveragePo	(None, 14, 14, 6)	0
10			
11	conv2d_1 (Conv2D)	(None, 10, 10, 16)	2416
12			
13	average_pooling2d_1 (Average	(None, 5, 5, 16)	0
14			
15	flatten (Flatten)	(None, 400)	0
16			
17	dense (Dense)	(None, 120)	48120
18			
19	dense_1 (Dense)	(None, 84)	10164
20			
21	dense_2 (Dense)	(None, 10)	850
22			
23	Total params: 61,706		
24	Trainable params: 61,706		
25	Non-trainable params: 0		
26			

2. Classical fully connected neural network implementation

```

1  def dense(input_shape, n_classes, lr):
2      model = tf.keras.models.Sequential([
3          tf.keras.layers.Flatten(input_shape=input_shape),
4          tf.keras.layers.Dense(128, activation='relu'),
5          tf.keras.layers.Dense(n_classes)
6      ])
7      model.compile(
8          optimizer=tf.keras.optimizers.Adam(lr),
9
10         loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
11         metrics=[tf.keras.metrics.SparseCategoricalAccuracy()],
12     )
13     return model

```

Model summary:

```

1  Model: "sequential"
2
3  _____
4  | Layer (type)                | Output Shape | Param # |
5  |-----|
6  | flatten_1 (Flatten)         | (None, 784)  | 0        |
7  |-----|
8  | dense_3 (Dense)             | (None, 128)  | 100480   |
9  |-----|
10 | dense_4 (Dense)             | (None, 10)   | 1290     |
11 |-----|
12 | Total params: 101,770
13 | Trainable params: 101,770
14 | Non-trainable params: 0

```

We used two different kinds of **keras** / **tensorflow** grammar to construct our neural network

As you can see the number of parameters differ greatly

Testing and training are all done with both of the model type

We used **tensorflow-datasets** to load the MNIST dataset, some enhancement are also done to it:

```

1  def load_mnist(batch_size=256):
2      (ds_train, ds_test, ds_val), ds_info = tfds.load(
3          'mnist',
4          split=[
5              tfds.Split.TRAIN,
6              tfds.Split.TEST.subsplit(tfds.percent[:50]),
7              tfds.Split.TEST.subsplit(tfds.percent[50:]),
8          ],
9          shuffle_files=True,
10         as_supervised=True,
11         with_info=True,
12     )

```

```

13
14     def normalize_img(image, label):
15         """Normalizes images: `uint8` → `float32`."""
16         return tf.cast(image, tf.float32) / 255., label
17
18     # Applying normalization before `ds.cache()` to re-use it.
19     # Note: Random transformations (e.g. images augmentations) should be
    applied
20     # after both `ds.cache()` (to avoid caching randomness) and `ds.batch()`
    (for
21     # vectorization [1]).
22     ds_train = ds_train.map(normalize_img,
    num_parallel_calls=tf.data.experimental.AUTOTUNE)
23     ds_train = ds_train.cache()
24     # For true randomness, we set the shuffle buffer to the full dataset size.
25     ds_train = ds_train.shuffle(ds_info.splits['train'].num_examples)
26     # Batch after shuffling to get unique batches at each epoch.
27     ds_train = ds_train.batch(batch_size)
28     ds_train = ds_train.prefetch(tf.data.experimental.AUTOTUNE)
29
30     ds_test = ds_test.map(normalize_img,
    num_parallel_calls=tf.data.experimental.AUTOTUNE)
31     ds_test = ds_test.batch(batch_size)
32     ds_test = ds_test.cache()
33     ds_test = ds_test.prefetch(tf.data.experimental.AUTOTUNE)
34
35     ds_val = ds_val.map(normalize_img,
    num_parallel_calls=tf.data.experimental.AUTOTUNE)
36     ds_val = ds_val.batch(batch_size)
37     ds_val = ds_val.cache()
38     ds_val = ds_val.prefetch(tf.data.experimental.AUTOTUNE)
39
40     return ds_train, ds_test, ds_val, ds_info

```

We split the **test** set provided by MNIST more thoroughly into **5000** validation images and **5000** test images.

Validations are to be used in training to save the best model (not used in this example)

(Gradient descent is done on training set, but the model with best validation accuracy are to be saved)

Tests are wild images that the network has never seen.

The main training logic would be:

```

1     def main(use_dense=False, index=0):
2         input_shape = (28, 28, 1)
3         n_classes = 10
4         batch_size = 256

```

```

5     lr = 1e-3
6     n_epoch = 20
7
8     result_dir = "results"
9     modelname = f"{result_dir}/models/{'dense' if use_dense else
'lenet5'}_{index}.pth"
10    figname = f"{result_dir}/logs/{'dense' if use_dense else
'lenet5'}_{index}.svg"
11    prediction_figname = f"prediction_{'dense' if use_dense else
'lenet5'}_{index}.svg"
12
13    log.info("Loading MNIST dataset")
14    ds_train, ds_test, ds_val, ds_info = load_mnist(batch_size)
15
16    log.info("Constructing LeNet-5")
17
18    construct_model = dense if use_dense else lenet5
19    model = construct_model(input_shape, n_classes, lr)
20    model.summary()
21
22    log.info("Training ... ")
23    history = model.fit(
24        ds_train,
25        epochs=n_epoch,
26        validation_data=ds_val,
27    )
28
29    log.info("Plotting training history ... ")
30    try:
31        plot_training_history(history, figname)
32    except Exception as e:
33        log.error("Cannot plot the training history (are you on GUI?), however
the figure is still saved to results/logs/{}_{}".format(index, modelname))
34
35    log.info("Evaluating ... ")
36    model.evaluate(ds_test)
37
38    log.info("Saving model ... ")
39
40    save_model(model, modelname)
41
42    log.info("Predicting on random images ... ")
43    predict_mnist(model, ds_test, prediction_figname)
44
45    return model, history, ds_train, ds_test, ds_val, ds_info

```

lr: learning rate, this is a hyperparameter. We found **1e-4** too slow

batch_size: this is also a hyperparameter. Training set of MNIST contains 60000 images of **28*28**. We can set the **batch_size** to 60000 if we want because this 2080ti do have the ability to handle that kind of input (consuming roughly 9000MB GPU Memory). But we such a huge batch size the learning gets quite slow across epochs. Probably due to the fact that hand-written number varies

just a little bit

In practice, we found **256** a fairly fast (in terms of both epoch time and loss descend) **batch_size**. **n_epoch** is the number of iterations we should run on the whole training set

In practice, one epoch of with a batch size of 256 would yield 85% accuracy across training set and validation set. The second epoch would yield nearly 95% of accuracy.

To plot the training history, we provide:

```
1  def plot_training_history(res, figname, limit_acc_tick=False):
2      def plot_key(key):
3          length = len(res.history[key])
4          inter = (interpolate.CubicSpline(np.linspace(0, length, length,
5          endpoint=False), res.history[key]))(np.linspace(0, length, length*10,
6          endpoint=False))
7          plt.plot(inter, label=key, linewidth=2.5, alpha=0.7)
8
9      plt.figure(figsize=(20, 10))
10     plt.suptitle("Training History: loss & val_loss & acc & val_acc",
11     fontweight="bold")
12     plt.subplot(121)
13     key = 'loss'
14     plot_key(key)
15     key = 'val_loss'
16     plot_key(key)
17     plt.legend(loc='right')
18
19     plt.subplot(122)
20     key = 'sparse_categorical_accuracy'
21     plot_key(key)
22     key = 'val_sparse_categorical_accuracy'
23     plot_key(key)
24     plt.legend(loc='right')
25     if limit_acc_tick:
26         plt.yticks(np.linspace(0, 1, 11, endpoint=True))
27
28     plt.savefig(figname)
29     plt.show()
```

Note that, we typically run this code without a GUI (**GNOME** or **Jupyter Notebook**), so we'd also save the figure to disk before trying to show it.

Similarly, we'd also provided a functions to do some prediction:


```

1  def predict_mnist(model, ds_test, figname='prediction.svg'):
2      try:
3          ds_test = ds_test.unbatch()
4          plt.figure(figsize=(10, 10))
5          plt.suptitle("Prediction & Ground Truth", fontweight="bold")
6          for i, (img, label) in enumerate(ds_test.take(9)):
7              plt.subplot(33*10 + i + 1)
8              img = np.expand_dims(img, 0)
9              result = model.predict(img)
10             result = np.argmax(result)
11             log.info(f"img shape: {img.shape}, label: {label}. predicted:
{result}")
12             plt.imshow(img.squeeze())
13             plt.title(f"Prediction Index: {i}, Prediction: {result}, Truth:
{label}")
14
15         plt.tight_layout()
16         plt.savefig(figname)
17         plt.show()

```

Results

A full training loop:

```

1  2021-01-16 13:25:32.766233: I
tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
opened dynamic library libcuda.so.1
2  2021-01-16 13:25:32.795200: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:981] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero
3  2021-01-16 13:25:32.795525: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1561] Found device 0 with
properties:
4  pciBusID: 0000:01:00.0 name: GeForce RTX 2080 Ti computeCapability: 7.5
5  coreClock: 1.545GHz coreCount: 68 deviceMemorySize: 10.75GiB
deviceMemoryBandwidth: 573.69GiB/s
6  2021-01-16 13:25:32.795674: I
tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
opened dynamic library libcudart.so.10.1
7  2021-01-16 13:25:32.796654: I
tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
opened dynamic library libcublas.so.10
8  2021-01-16 13:25:32.797596: I
tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
opened dynamic library libcufft.so.10

```

9 2021-01-16 13:25:32.797765: I
tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
opened dynamic library libcurand.so.10

10 2021-01-16 13:25:32.798743: I
tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
opened dynamic library libcusolver.so.10

11 2021-01-16 13:25:32.799294: I
tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
opened dynamic library libcusparsparse.so.10

12 2021-01-16 13:25:32.801423: I
tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
opened dynamic library libcudnn.so.7

13 2021-01-16 13:25:32.801535: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:981] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero

14 2021-01-16 13:25:32.801890: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:981] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero

15 2021-01-16 13:25:32.802180: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1703] Adding visible gpu
devices: 0

16 2021-01-16 13:25:32.802417: I
tensorflow/core/platform/cpu_feature_guard.cc:143] Your CPU supports
instructions that this TensorFlow binary was not compiled to use: SSE4.1
SSE4.2 AVX AVX2 FMA

17 2021-01-16 13:25:32.806249: I
tensorflow/core/platform/profile_utils/cpu_utils.cc:102] CPU Frequency:
3000000000 Hz

18 2021-01-16 13:25:32.806487: I tensorflow/compiler/xla/service/service.cc:168]
XLA service 0x5652d08e4fc0 initialized for platform Host (this does not
guarantee that XLA will be used). Devices:

19 2021-01-16 13:25:32.806520: I tensorflow/compiler/xla/service/service.cc:176]
StreamExecutor device (0): Host, Default Version

20 2021-01-16 13:25:32.806654: I
tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:981] successful NUMA node
read from SysFS had negative value (-1), but there must be at least one NUMA
node, so returning NUMA node zero

21 2021-01-16 13:25:32.806987: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1561] Found device 0 with
properties:

22 pciBusID: 0000:01:00.0 name: GeForce RTX 2080 Ti computeCapability: 7.5

23 coreClock: 1.545GHz coreCount: 68 deviceMemorySize: 10.75GiB

24 deviceMemoryBandwidth: 573.69GiB/s

25 2021-01-16 13:25:32.807016: I
tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
opened dynamic library libcudart.so.10.1

26 2021-01-16 13:25:32.807044: I
tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
opened dynamic library libcublas.so.10

```
26 2021-01-16 13:25:32.807068: I
    tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
    opened dynamic library libcufft.so.10
27 2021-01-16 13:25:32.807093: I
    tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
    opened dynamic library libcurand.so.10
28 2021-01-16 13:25:32.807104: I
    tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
    opened dynamic library libcusolver.so.10
29 2021-01-16 13:25:32.807114: I
    tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
    opened dynamic library libcusparsesparse.so.10
30 2021-01-16 13:25:32.807138: I
    tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
    opened dynamic library libcudnn.so.7
31 2021-01-16 13:25:32.807187: I
    tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:981] successful NUMA node
    read from SysFS had negative value (-1), but there must be at least one NUMA
    node, so returning NUMA node zero
32 2021-01-16 13:25:32.807486: I
    tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:981] successful NUMA node
    read from SysFS had negative value (-1), but there must be at least one NUMA
    node, so returning NUMA node zero
33 2021-01-16 13:25:32.807791: I
    tensorflow/core/common_runtime/gpu/gpu_device.cc:1703] Adding visible gpu
    devices: 0
34 2021-01-16 13:25:32.807825: I
    tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
    opened dynamic library libcudart.so.10.1
35 2021-01-16 13:25:32.874355: I
    tensorflow/core/common_runtime/gpu/gpu_device.cc:1102] Device interconnect
    StreamExecutor with strength 1 edge matrix:
36 2021-01-16 13:25:32.874390: I
    tensorflow/core/common_runtime/gpu/gpu_device.cc:1108]          0
37 2021-01-16 13:25:32.874395: I
    tensorflow/core/common_runtime/gpu/gpu_device.cc:1121] 0:    N
38 2021-01-16 13:25:32.874556: I
    tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:981] successful NUMA node
    read from SysFS had negative value (-1), but there must be at least one NUMA
    node, so returning NUMA node zero
39 2021-01-16 13:25:32.874901: I
    tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:981] successful NUMA node
    read from SysFS had negative value (-1), but there must be at least one NUMA
    node, so returning NUMA node zero
40 2021-01-16 13:25:32.875213: I
    tensorflow/stream_executor/cuda/cuda_gpu_executor.cc:981] successful NUMA node
    read from SysFS had negative value (-1), but there must be at least one NUMA
    node, so returning NUMA node zero
```

```

41 2021-01-16 13:25:32.875508: I
tensorflow/core/common_runtime/gpu/gpu_device.cc:1247] Created TensorFlow
device (/job:localhost/replica:0/task:0/device:GPU:0 with 10094 MB memory) →
physical GPU (device: 0, name: GeForce RTX 2080 Ti, pci bus id: 0000:01:00.0,
compute capability: 7.5)
42 2021-01-16 13:25:32.876812: I tensorflow/compiler/xla/service/service.cc:168]
XLA service 0x5652d42dc050 initialized for platform CUDA (this does not
guarantee that XLA will be used). Devices:
43 2021-01-16 13:25:32.876822: I tensorflow/compiler/xla/service/service.cc:176]
StreamExecutor device (0): GeForce RTX 2080 Ti, Compute Capability 7.5
44 1 Physical GPUs, 1 Logical GPUs
45 2021-01-16 13:25:32 xzdd-ubuntu __main__[167776] INFO Loading MNIST dataset
46 2021-01-16 13:25:32 xzdd-ubuntu absl[167776] INFO Overwrite dataset info from
restored data version.
47 2021-01-16 13:25:32 xzdd-ubuntu absl[167776] INFO Reusing dataset mnist
(/home/xzdd/tensorflow_datasets/mnist/1.0.0)
48 2021-01-16 13:25:32 xzdd-ubuntu absl[167776] INFO Constructing tf.data.Dataset
for split [NamedSplit('train'), NamedSplit('test')(tfds.percent[:50]),
NamedSplit('test')(tfds.percent[50:])], from
/home/xzdd/tensorflow_datasets/mnist/1.0.0
49 2021-01-16 13:25:33 xzdd-ubuntu __main__[167776] INFO Constructing LeNet-5
50 Model: "model"
51
52 Layer (type) Output Shape Param #
53
54 input_1 (InputLayer) [(None, 28, 28, 1)] 0
55
56 conv2d (Conv2D) (None, 28, 28, 6) 156
57
58 average_pooling2d (AveragePo (None, 14, 14, 6) 0
59
60 conv2d_1 (Conv2D) (None, 10, 10, 16) 2416
61
62 average_pooling2d_1 (Average (None, 5, 5, 16) 0
63
64 flatten (Flatten) (None, 400) 0
65
66 dense (Dense) (None, 120) 48120
67
68 dense_1 (Dense) (None, 84) 10164
69
70 dense_2 (Dense) (None, 10) 850
71
72 Total params: 61,706
73 Trainable params: 61,706
74 Non-trainable params: 0
75
76 2021-01-16 13:25:33 xzdd-ubuntu __main__[167776] INFO Training...
77 Epoch 1/20
78 2021-01-16 13:25:33.610233: I
tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
opened dynamic library libcublas.so.10

```

```
79 2021-01-16 13:25:34.677403: I
    tensorflow/stream_executor/platform/default/dso_loader.cc:44] Successfully
    opened dynamic library libcudnn.so.7
80 2021-01-16 13:25:35.275715: W
    tensorflow/stream_executor/gpu/asm_compiler.cc:81] Running ptxas --version
    returned 256
81 2021-01-16 13:25:35.304668: W
    tensorflow/stream_executor/gpu/redzone_allocator.cc:314] Internal: ptxas
    exited with non-zero error code 256, output:
82 Relying on driver to perform ptx compilation.
83 Modify $PATH to customize ptxas location.
84 This message will be only logged once.
85 235/235 [=====] - 6s 26ms/step - loss: 1.6721 -
    sparse_categorical_accuracy: 0.8184 - val_loss: 1.5410 -
    val_sparse_categorical_accuracy: 0.9308
86 Epoch 2/20
87 235/235 [=====] - 1s 4ms/step - loss: 1.5264 -
    sparse_categorical_accuracy: 0.9432 - val_loss: 1.5089 -
    val_sparse_categorical_accuracy: 0.9604
88 Epoch 3/20
89 235/235 [=====] - 2s 7ms/step - loss: 1.5047 -
    sparse_categorical_accuracy: 0.9622 - val_loss: 1.4958 -
    val_sparse_categorical_accuracy: 0.9700
90 Epoch 4/20
91 235/235 [=====] - 1s 5ms/step - loss: 1.4943 -
    sparse_categorical_accuracy: 0.9712 - val_loss: 1.4928 -
    val_sparse_categorical_accuracy: 0.9712
92 Epoch 5/20
93 235/235 [=====] - 1s 3ms/step - loss: 1.4881 -
    sparse_categorical_accuracy: 0.9766 - val_loss: 1.4860 -
    val_sparse_categorical_accuracy: 0.9782
94 Epoch 6/20
95 235/235 [=====] - 1s 5ms/step - loss: 1.4838 -
    sparse_categorical_accuracy: 0.9805 - val_loss: 1.4844 -
    val_sparse_categorical_accuracy: 0.9798
96 Epoch 7/20
97 235/235 [=====] - 1s 5ms/step - loss: 1.4806 -
    sparse_categorical_accuracy: 0.9830 - val_loss: 1.4809 -
    val_sparse_categorical_accuracy: 0.9828
98 Epoch 8/20
99 235/235 [=====] - 2s 7ms/step - loss: 1.4784 -
    sparse_categorical_accuracy: 0.9850 - val_loss: 1.4804 -
    val_sparse_categorical_accuracy: 0.9822
100 Epoch 9/20
101 235/235 [=====] - 2s 7ms/step - loss: 1.4763 -
    sparse_categorical_accuracy: 0.9868 - val_loss: 1.4802 -
    val_sparse_categorical_accuracy: 0.9828
102 Epoch 10/20
103 235/235 [=====] - 1s 4ms/step - loss: 1.4748 -
    sparse_categorical_accuracy: 0.9879 - val_loss: 1.4807 -
    val_sparse_categorical_accuracy: 0.9814
104 Epoch 11/20
```

```
105 235/235 [=====] - 2s 6ms/step - loss: 1.4735 -  
    sparse_categorical_accuracy: 0.9892 - val_loss: 1.4784 -  
    val_sparse_categorical_accuracy: 0.9840  
106 Epoch 12/20  
107 235/235 [=====] - 2s 7ms/step - loss: 1.4725 -  
    sparse_categorical_accuracy: 0.9900 - val_loss: 1.4794 -  
    val_sparse_categorical_accuracy: 0.9832  
108 Epoch 13/20  
109 235/235 [=====] - 1s 3ms/step - loss: 1.4711 -  
    sparse_categorical_accuracy: 0.9916 - val_loss: 1.4776 -  
    val_sparse_categorical_accuracy: 0.9848  
110 Epoch 14/20  
111 235/235 [=====] - 1s 2ms/step - loss: 1.4707 -  
    sparse_categorical_accuracy: 0.9916 - val_loss: 1.4781 -  
    val_sparse_categorical_accuracy: 0.9848  
112 Epoch 15/20  
113 235/235 [=====] - 1s 3ms/step - loss: 1.4698 -  
    sparse_categorical_accuracy: 0.9925 - val_loss: 1.4776 -  
    val_sparse_categorical_accuracy: 0.9848  
114 Epoch 16/20  
115 235/235 [=====] - 1s 3ms/step - loss: 1.4694 -  
    sparse_categorical_accuracy: 0.9927 - val_loss: 1.4780 -  
    val_sparse_categorical_accuracy: 0.9834  
116 Epoch 17/20  
117 235/235 [=====] - 1s 6ms/step - loss: 1.4685 -  
    sparse_categorical_accuracy: 0.9935 - val_loss: 1.4771 -  
    val_sparse_categorical_accuracy: 0.9842  
118 Epoch 18/20  
119 235/235 [=====] - 1s 5ms/step - loss: 1.4683 -  
    sparse_categorical_accuracy: 0.9936 - val_loss: 1.4777 -  
    val_sparse_categorical_accuracy: 0.9846  
120 Epoch 19/20  
121 235/235 [=====] - 1s 2ms/step - loss: 1.4672 -  
    sparse_categorical_accuracy: 0.9946 - val_loss: 1.4757 -  
    val_sparse_categorical_accuracy: 0.9866  
122 Epoch 20/20  
123 235/235 [=====] - 1s 4ms/step - loss: 1.4671 -  
    sparse_categorical_accuracy: 0.9949 - val_loss: 1.4764 -  
    val_sparse_categorical_accuracy: 0.9858  
124 2021-01-16 13:26:05 xzdd-ubuntu __main__[167776] INFO Plotting training  
    history ...  
125 2021-01-16 13:26:05 xzdd-ubuntu __main__[167776] INFO Evaluating ...  
126 20/20 [=====] - 0s 7ms/step - loss: 1.4758 -  
    sparse_categorical_accuracy: 0.9866  
127 2021-01-16 13:26:05 xzdd-ubuntu __main__[167776] INFO Saving model ...  
128 2021-01-16 13:26:06.153277: W tensorflow/python/util/util.cc:329] Sets are not  
    currently considered sequences, but this may change in the future, so consider  
    avoiding using them.
```

```

129 WARNING:tensorflow:From /home/xzdd/miniconda3/envs/py38/lib/python3.8/site-
    packages/tensorflow/python/ops/resource_variable_ops.py:1813: calling
    BaseResourceVariable.__init__ (from
    tensorflow.python.ops.resource_variable_ops) with constraint is deprecated and
    will be removed in a future version.
130 Instructions for updating:
131 If using Keras pass *_constraint arguments to layers.
132 2021-01-16 13:26:06 xzdd-ubuntu tensorflow[167776] WARNING From
    /home/xzdd/miniconda3/envs/py38/lib/python3.8/site-
    packages/tensorflow/python/ops/resource_variable_ops.py:1813: calling
    BaseResourceVariable.__init__ (from
    tensorflow.python.ops.resource_variable_ops) with constraint is deprecated and
    will be removed in a future version.
133 Instructions for updating:
134 If using Keras pass *_constraint arguments to layers.
135 INFO:tensorflow:Assets written to: results/models/lenet5_0.pth/assets
136 2021-01-16 13:26:06 xzdd-ubuntu tensorflow[167776] INFO Assets written to:
    results/models/lenet5_0.pth/assets
137 2021-01-16 13:26:06 xzdd-ubuntu __main__[167776] INFO Predicting on random
    images ...
138 2021-01-16 13:26:08 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 3. predicted: 3
139 2021-01-16 13:26:08 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 6. predicted: 6
140 2021-01-16 13:26:08 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 7. predicted: 7
141 2021-01-16 13:26:08 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 2. predicted: 2
142 2021-01-16 13:26:08 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 3. predicted: 3
143 2021-01-16 13:26:08 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 2. predicted: 2
144 2021-01-16 13:26:08 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 4. predicted: 4
145 2021-01-16 13:26:08 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 7. predicted: 7
146 2021-01-16 13:26:08 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 9. predicted: 9
147 2021-01-16 13:26:08 xzdd-ubuntu __main__[167776] INFO Loading MNIST dataset
148 2021-01-16 13:26:08 xzdd-ubuntu absl[167776] INFO Overwrite dataset info from
    restored data version.
149 2021-01-16 13:26:08 xzdd-ubuntu absl[167776] INFO Reusing dataset mnist
    (/home/xzdd/tensorflow_datasets/mnist/1.0.0)
150 2021-01-16 13:26:08 xzdd-ubuntu absl[167776] INFO Constructing tf.data.Dataset
    for split [NamedSplit('train'), NamedSplit('test')(tfds.percent[:50]),
    NamedSplit('test')(tfds.percent[50:])] from
    /home/xzdd/tensorflow_datasets/mnist/1.0.0
151 2021-01-16 13:26:08 xzdd-ubuntu __main__[167776] INFO Constructing LeNet-5
152 Model: "sequential"
153
154 Layer (type)                Output Shape                Param #
155

```

```

156 flatten_1 (Flatten)          (None, 784)          0
157
158 dense_3 (Dense)              (None, 128)         100480
159
160 dense_4 (Dense)              (None, 10)          1290
161
162 Total params: 101,770
163 Trainable params: 101,770
164 Non-trainable params: 0
165
166 2021-01-16 13:26:08 xzdd-ubuntu __main__[167776] INFO Training...
167 Epoch 1/20
168 235/235 [=====] - 0s 2ms/step - loss: 0.4406 -
  sparse_categorical_accuracy: 0.8840 - val_loss: 0.2366 -
  val_sparse_categorical_accuracy: 0.9324
169 Epoch 2/20
170 235/235 [=====] - 0s 2ms/step - loss: 0.2010 -
  sparse_categorical_accuracy: 0.9434 - val_loss: 0.1740 -
  val_sparse_categorical_accuracy: 0.9476
171 Epoch 3/20
172 235/235 [=====] - 0s 1ms/step - loss: 0.1501 -
  sparse_categorical_accuracy: 0.9571 - val_loss: 0.1396 -
  val_sparse_categorical_accuracy: 0.9596
173 Epoch 4/20
174 235/235 [=====] - 0s 1ms/step - loss: 0.1194 -
  sparse_categorical_accuracy: 0.9662 - val_loss: 0.1206 -
  val_sparse_categorical_accuracy: 0.9630
175 Epoch 5/20
176 235/235 [=====] - 0s 1ms/step - loss: 0.0991 -
  sparse_categorical_accuracy: 0.9714 - val_loss: 0.1072 -
  val_sparse_categorical_accuracy: 0.9666
177 Epoch 6/20
178 235/235 [=====] - 0s 1ms/step - loss: 0.0840 -
  sparse_categorical_accuracy: 0.9760 - val_loss: 0.0967 -
  val_sparse_categorical_accuracy: 0.9700
179 Epoch 7/20
180 235/235 [=====] - 0s 1ms/step - loss: 0.0723 -
  sparse_categorical_accuracy: 0.9798 - val_loss: 0.0919 -
  val_sparse_categorical_accuracy: 0.9708
181 Epoch 8/20
182 235/235 [=====] - 0s 1ms/step - loss: 0.0626 -
  sparse_categorical_accuracy: 0.9826 - val_loss: 0.0810 -
  val_sparse_categorical_accuracy: 0.9756
183 Epoch 9/20
184 235/235 [=====] - 0s 1ms/step - loss: 0.0548 -
  sparse_categorical_accuracy: 0.9847 - val_loss: 0.0816 -
  val_sparse_categorical_accuracy: 0.9750
185 Epoch 10/20
186 235/235 [=====] - 0s 1ms/step - loss: 0.0491 -
  sparse_categorical_accuracy: 0.9861 - val_loss: 0.0771 -
  val_sparse_categorical_accuracy: 0.9764
187 Epoch 11/20

```



```
188 235/235 [=====] - 0s 1ms/step - loss: 0.0425 -  
    sparse_categorical_accuracy: 0.9882 - val_loss: 0.0727 -  
    val_sparse_categorical_accuracy: 0.9770  
189 Epoch 12/20  
190 235/235 [=====] - 0s 2ms/step - loss: 0.0374 -  
    sparse_categorical_accuracy: 0.9900 - val_loss: 0.0715 -  
    val_sparse_categorical_accuracy: 0.9784  
191 Epoch 13/20  
192 235/235 [=====] - 0s 1ms/step - loss: 0.0333 -  
    sparse_categorical_accuracy: 0.9916 - val_loss: 0.0690 -  
    val_sparse_categorical_accuracy: 0.9778  
193 Epoch 14/20  
194 235/235 [=====] - 0s 1ms/step - loss: 0.0296 -  
    sparse_categorical_accuracy: 0.9925 - val_loss: 0.0711 -  
    val_sparse_categorical_accuracy: 0.9782  
195 Epoch 15/20  
196 235/235 [=====] - 0s 1ms/step - loss: 0.0261 -  
    sparse_categorical_accuracy: 0.9938 - val_loss: 0.0700 -  
    val_sparse_categorical_accuracy: 0.9762  
197 Epoch 16/20  
198 235/235 [=====] - 0s 2ms/step - loss: 0.0237 -  
    sparse_categorical_accuracy: 0.9948 - val_loss: 0.0676 -  
    val_sparse_categorical_accuracy: 0.9784  
199 Epoch 17/20  
200 235/235 [=====] - 0s 1ms/step - loss: 0.0204 -  
    sparse_categorical_accuracy: 0.9957 - val_loss: 0.0689 -  
    val_sparse_categorical_accuracy: 0.9778  
201 Epoch 18/20  
202 235/235 [=====] - 0s 1ms/step - loss: 0.0187 -  
    sparse_categorical_accuracy: 0.9960 - val_loss: 0.0686 -  
    val_sparse_categorical_accuracy: 0.9780  
203 Epoch 19/20  
204 235/235 [=====] - 0s 2ms/step - loss: 0.0166 -  
    sparse_categorical_accuracy: 0.9968 - val_loss: 0.0694 -  
    val_sparse_categorical_accuracy: 0.9788  
205 Epoch 20/20  
206 235/235 [=====] - 0s 2ms/step - loss: 0.0145 -  
    sparse_categorical_accuracy: 0.9973 - val_loss: 0.0660 -  
    val_sparse_categorical_accuracy: 0.9792  
207 2021-01-16 13:26:18 xzdd-ubuntu __main__[167776] INFO Plotting training  
    history ...  
208 2021-01-16 13:26:18 xzdd-ubuntu __main__[167776] INFO Evaluating ...  
209 20/20 [=====] - 0s 7ms/step - loss: 0.0736 -  
    sparse_categorical_accuracy: 0.9780  
210 2021-01-16 13:26:18 xzdd-ubuntu __main__[167776] INFO Saving model ...  
211 INFO:tensorflow:Assets written to: results/models/dense_0.pth/assets  
212 2021-01-16 13:26:18 xzdd-ubuntu tensorflow[167776] INFO Assets written to:  
    results/models/dense_0.pth/assets  
213 2021-01-16 13:26:18 xzdd-ubuntu __main__[167776] INFO Predicting on random  
    images ...  
214 2021-01-16 13:26:18 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,  
    1), label: 3. predicted: 3
```


```

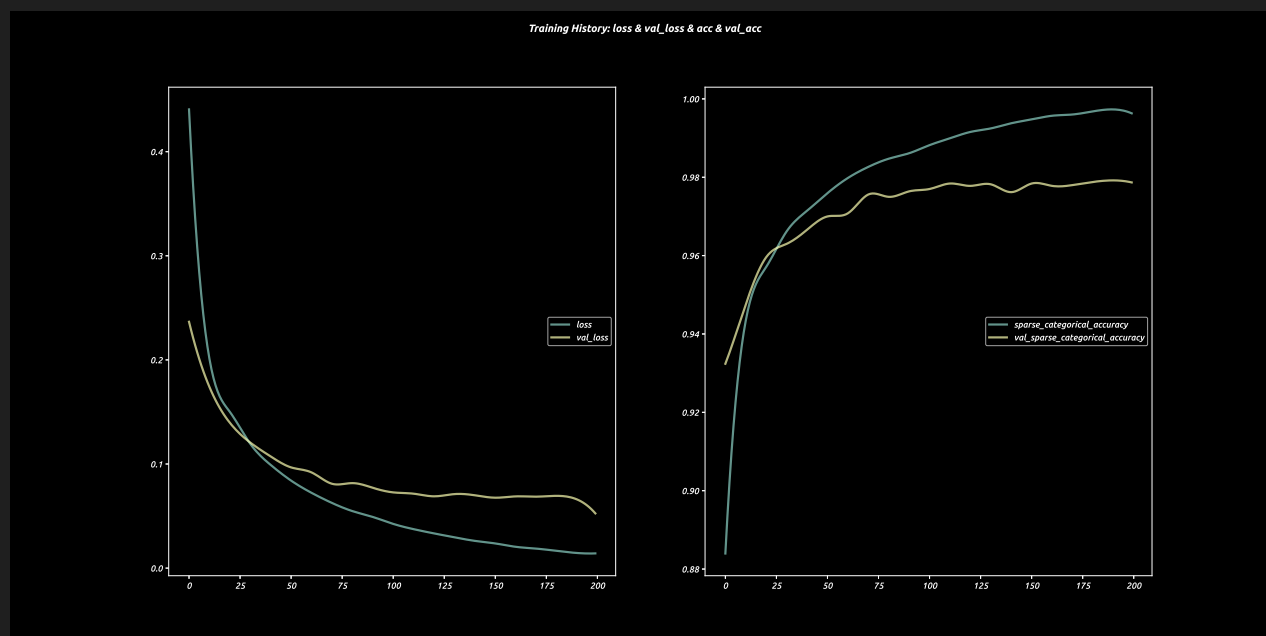
215 2021-01-16 13:26:18 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 7. predicted: 7
216 2021-01-16 13:26:18 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 6. predicted: 6
217 2021-01-16 13:26:18 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 2. predicted: 2
218 2021-01-16 13:26:18 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 2. predicted: 2
219 2021-01-16 13:26:18 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 2. predicted: 2
220 2021-01-16 13:26:19 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 3. predicted: 3
221 2021-01-16 13:26:19 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 4. predicted: 4
222 2021-01-16 13:26:19 xzdd-ubuntu __main__[167776] INFO img shape: (1, 28, 28,
    1), label: 7. predicted: 7

```

Figures in Training and Prediction sections are all of vector format, zoom in if you feel uncomfortable

Training





Evaluation

LeNet-5

```
2021-01-16 13:26:05 xzdd-ubuntu __main__[167776] INFO Evaluating...  
20/20 [=====] - 0s 7ms/step - loss: 1.4758 - sparse_categorical_accuracy: 0.9866
```


Dense

```
2021-01-16 13:26:18 xzdd-ubuntu __main__[167776] INFO Evaluating...  
20/20 [=====] - 0s 7ms/step - loss: 0.0736 - sparse_categorical_accuracy: 0.9780
```


Prediction

Prediction & Ground Truth


Prediction Index: 0, Prediction: 3, Truth: 3




Prediction Index: 1, Prediction: 6, Truth: 6




Prediction Index: 2, Prediction: 7, Truth: 7




Prediction Index: 3, Prediction: 2, Truth: 2




Prediction Index: 4, Prediction: 3, Truth: 3




Prediction Index: 5, Prediction: 2, Truth: 2




Prediction Index: 6, Prediction: 4, Truth: 4



Prediction Index: 7, Prediction: 7, Truth: 7




Prediction Index: 8, Prediction: 9, Truth: 9




Prediction & Ground Truth


Prediction Index: 0, Prediction: 3, Truth: 3




Prediction Index: 1, Prediction: 7, Truth: 7




Prediction Index: 2, Prediction: 6, Truth: 6




Prediction Index: 3, Prediction: 2, Truth: 2




Prediction Index: 4, Prediction: 2, Truth: 2



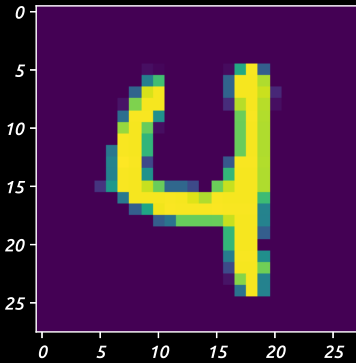
Prediction Index: 5, Prediction: 2, Truth: 2



Prediction Index: 6, Prediction: 3, Truth: 3



Prediction Index: 7, Prediction: 4, Truth: 4



Prediction Index: 8, Prediction: 7, Truth: 7

